

DIP_Assignment (/github/keithbored3310/DIP_Assignment/tree/main)
/ JupyterNotebook_G1.ipynb (/github/keithbored3310/DIP_Assignment/tree/main/JupyterNotebook_G1.ipynb)

Import Libraries

```
In [79]: # Import the necessary Libraries
import os
import numpy as np
import cv2
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix, precision_recall_curve
from sklearn.preprocessing import label_binarize
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers, regularizers
from tensorflow.keras.applications import ResNet50
import matplotlib.pyplot as plt
```

Load Dataset

```
In [80]: # Load the image data
image_dir = './dtd-r1.0.1/dtd/images/'
class_names = os.listdir(image_dir)

X = []
y = []

for i, class_name in enumerate(class_names):
    class_dir = os.path.join(image_dir, class_name)
    for image_path in os.listdir(class_dir):
        image = keras.preprocessing.image.load_img(os.path.join(class_dir, image_path))
        image_array = keras.preprocessing.image.img_to_array(image)
        X.append(image_array)
        y.append(i)

# Convert the lists to numpy arrays
X = np.array(X)
y = np.array(y)
```

```
<ipython-input-80-ddfe85779b19>:17: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray.
```

```
X = np.array(X)
```

Explore Dataset

```
In [86]: # Define the number of images to preview per class
num_images_per_class = 5

# Preview images for each class
for class_label, class_name in enumerate(class_names):
    # Find the indices of images belonging to the current class
    class_indices = np.where(y == class_label)[0]

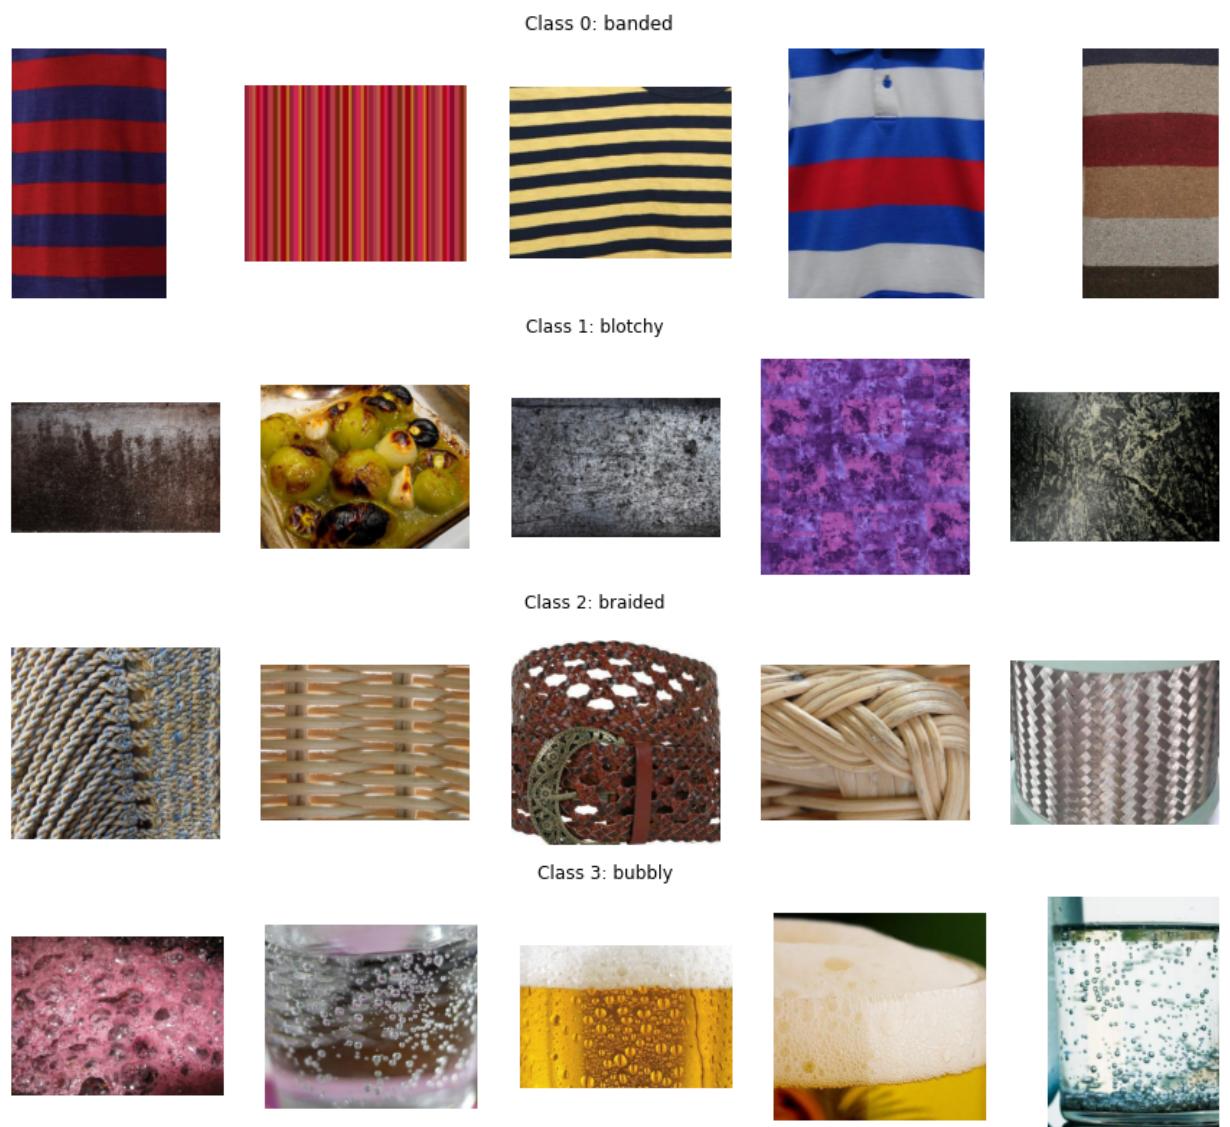
    # Check if there are images available for the current class
    if len(class_indices) == 0:
        print(f"No images available for Class {class_label}: {class_name}")
        continue

    # Randomly select a subset of images to preview
    preview_indices = np.random.choice(class_indices, size=min(num_images_per_class, len(class_))

    # Plot the preview images
    fig, axes = plt.subplots(1, len(preview_indices), figsize=(15, 3))
    for i, index in enumerate(preview_indices):
        image = X[index].astype(np.uint8)
        axes[i].imshow(image)
        axes[i].axis('off')

    # Set the title of the plot as the class name
    fig.suptitle(f'Class {class_label}: {class_name}')

    # Show the plot
    plt.show()
```



Class 4: bumpy



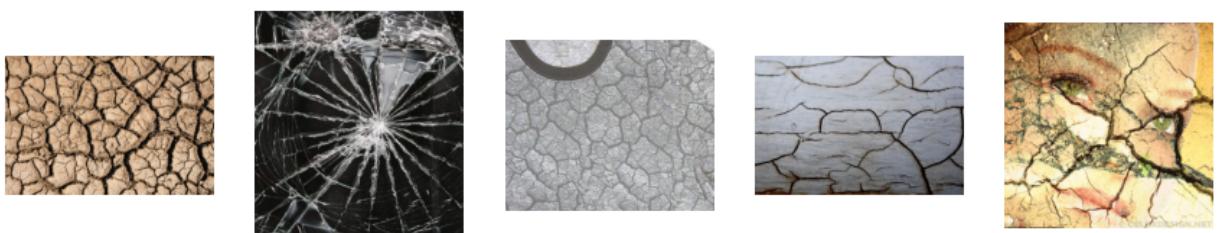
Class 5: chequered



Class 6: cobwebbed



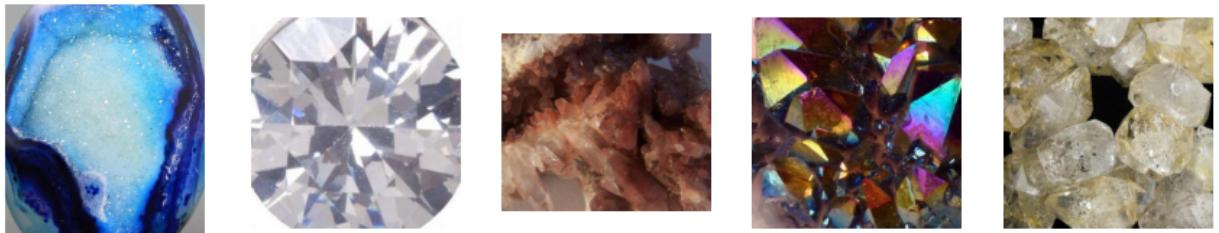
Class 7: cracked



Class 8: crosshatched



Class 9: crystalline



Class 10: dotted



Class 11: fibrous



Class 12: flecked



Class 13: freckled



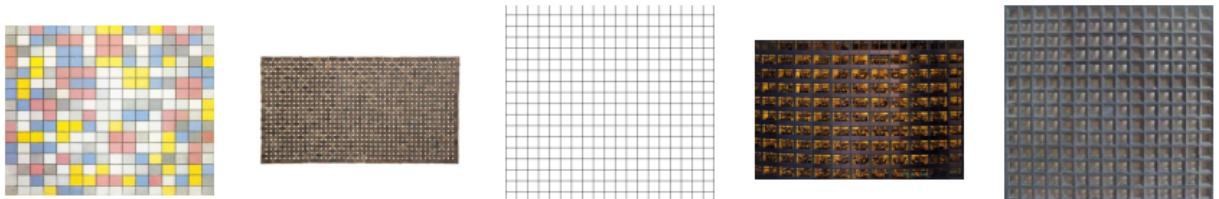
Class 14: frilly



Class 15: gauzy



Class 16: grid



Class 17: grooved



Class 18: honeycombed



Class 19: interlaced



Class 20: knitted



Class 21: lacelike



Class 22: lined



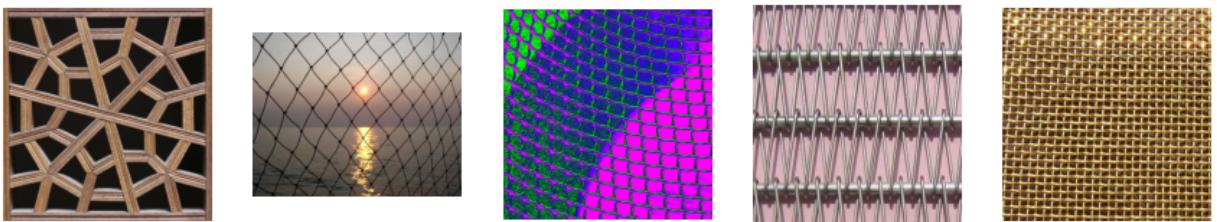
Class 23: marbled



Class 24: matted



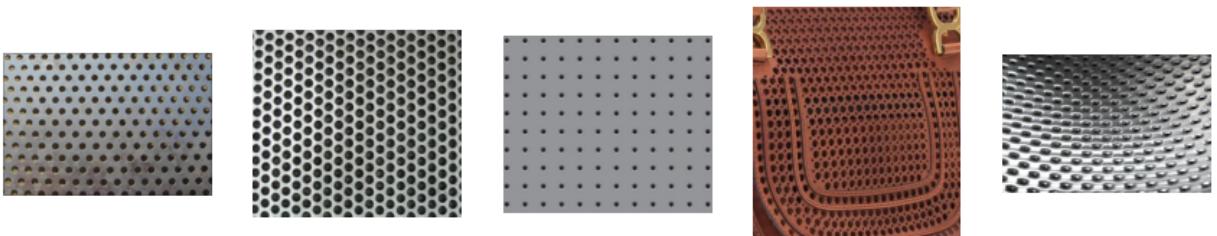
Class 25: meshed



Class 26: paisley



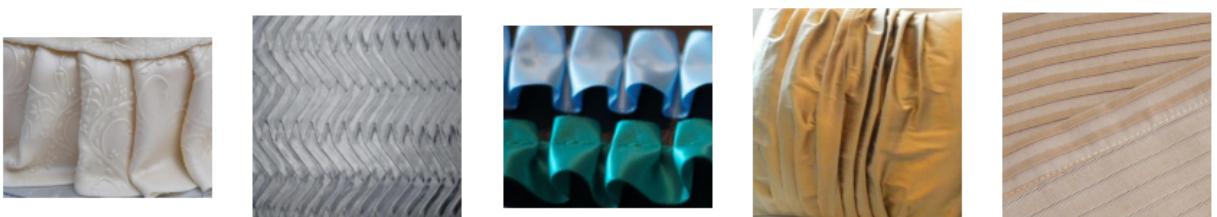
Class 27: perforated



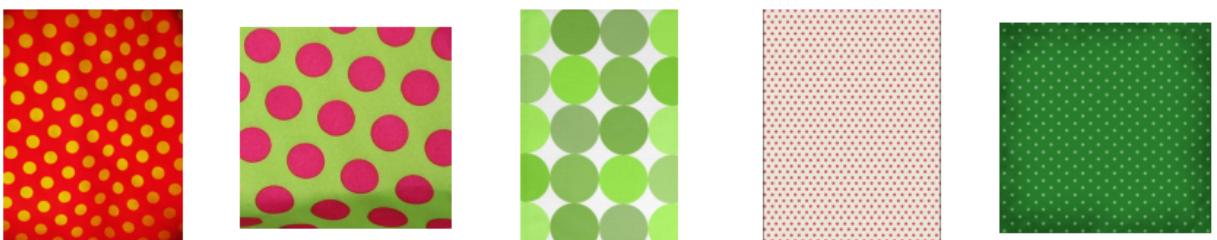
Class 28: pitted



Class 29: pleated



Class 30: polka-dotted



Class 31: porous



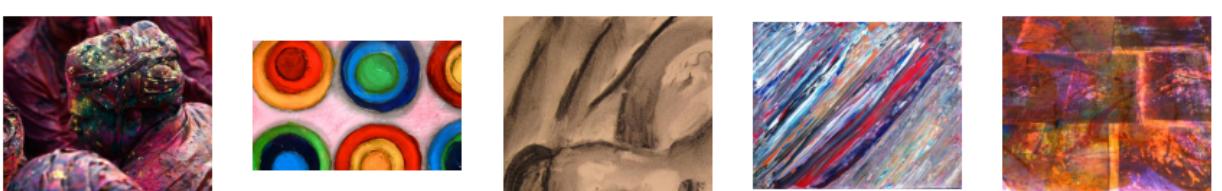
Class 32: potholed



Class 33: scaly



Class 34: smeared



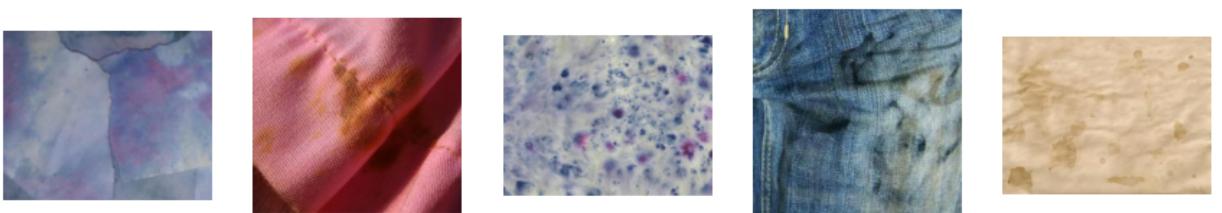
Class 35: spiralled



Class 36: sprinkled



Class 37: stained



Class 38: stratified



Class 39: striped



Class 40: studded



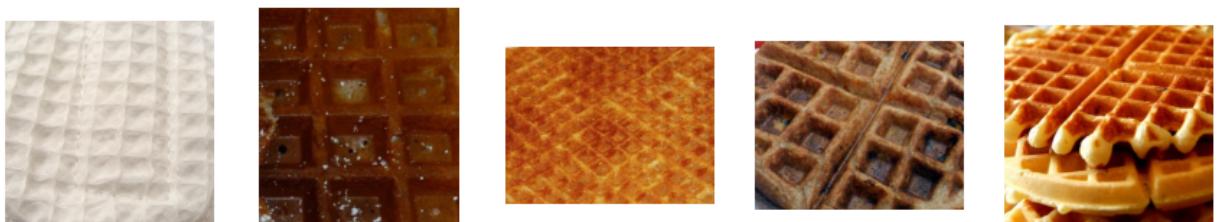
Class 41: swirly



Class 42: veined



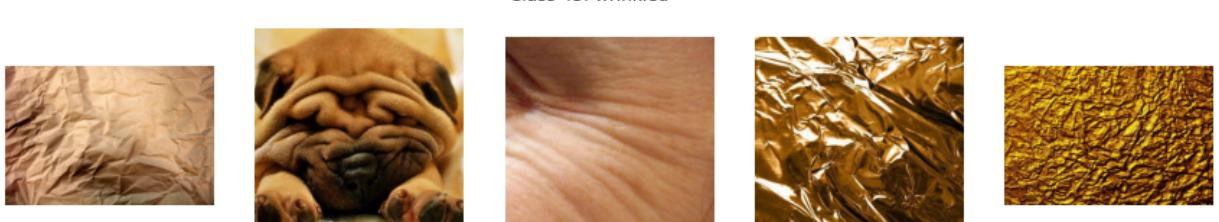
Class 43: waffled



Class 44: woven



Class 45: wrinkled



Class 46: zigzagged



```
In [87]: # Count the total number of images
total_images = len(X)
print("Total number of images:", total_images)

# Determine the image size and type
image_size = X[0].shape
image_type = type(X[1][0][0][0])
print("Image size:", image_size)
print("Image type:", image_type)
```

Total number of images: 5640
 Image size: (300, 300, 3)
 Image type: <class 'numpy.float32'>

```
In [95]: # Check the shape of the data
print("X shape:", X.shape)
print("y shape:", y.shape)

# Check the number of classes
num_classes = len(class_names)
print("Number of classes:", num_classes)

# Check the distribution of classes
class_counts = np.bincount(y)
print("Class counts:", class_counts)
```

X shape: (5640,)
 y shape: (5640,)
 Number of classes: 47
 Class counts: [120 120]

Split Dataset

```
In [89]: # Split the data into training and combined set (testing + validation)
X_train, X_combined, y_train, y_combined = train_test_split(X, y, test_size=0.2, random_state=42)

# Split the combined set into testing and validation sets
X_test, X_val, y_test, y_val = train_test_split(X_combined, y_combined, test_size=0.5, random_state=42)

# Now there are X_train, X_test, X_val as the respective feature sets
# and y_train, y_test, y_val as the respective target sets
```

```
In [90]: print("X_train shape:", X_train.shape)
print("X_val shape:", X_val.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_val shape:", y_val.shape)
print("y_test shape:", y_test.shape)
print("Number of samples in X_train:", len(X_train))
print("Number of samples in X_val:", len(X_val))
print("Number of samples in X_test:", len(X_test))
print("Number of labels in y_train:", len(y_train))
print("Number of labels in y_val:", len(y_val))
print("Number of labels in y_test:", len(y_test))
```

```
X_train shape: (4512,)
X_val shape: (564,)
X_test shape: (564,)
y_train shape: (4512,)
y_val shape: (564,)
y_test shape: (564,)
Number of samples in X_train: 4512
Number of samples in X_val: 564
Number of samples in X_test: 564
Number of labels in y_train: 4512
Number of labels in y_val: 564
Number of labels in y_test: 564
```

Pre-process Dataset

```
In [91]: resized_X_train = []
resized_X_val = []
resized_X_test = []

# Resize images in the training set
for image in X_train:
    resized_image = cv2.resize(image, (224, 224))
    resized_X_train.append(resized_image)

# Resize images in the validation set
for image in X_val:
    resized_image = cv2.resize(image, (224, 224))
    resized_X_val.append(resized_image)

# Resize images in the testing set
for image in X_test:
    resized_image = cv2.resize(image, (224, 224))
    resized_X_test.append(resized_image)

# Convert the lists to numpy arrays
X_train = np.array(resized_X_train)
X_val = np.array(resized_X_val)
X_test = np.array(resized_X_test)

# Reshape the input data to the desired shape
X_train = X_train.reshape(X_train.shape[0], 224, 224, 3)
X_val = X_val.reshape(X_val.shape[0], 224, 224, 3)
X_test = X_test.reshape(X_test.shape[0], 224, 224, 3)
```

```
In [92]: # Normalize X
X_train = X_train / 255.0
X_val = X_val / 255.0
X_test = X_test / 255.0
```

```
In [93]: # Perform one-hot encoding on the labels
y_train = keras.utils.to_categorical(y_train, num_classes)
y_val = keras.utils.to_categorical(y_val, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

LeNet 5 Model

```
In [10]: # Define the LeNet 5 model architecture
model = keras.Sequential([
    layers.Conv2D(6,kernel_size=(5,5),activation='relu',input_shape=(224,224,3)),
    layers.MaxPooling2D(pool_size=(2,2)),
    layers.Conv2D(16, kernel_size=(5,5), activation='relu'),
    layers.MaxPooling2D(pool_size=(2,2)),
    layers.Flatten(),
    layers.Dense(120, activation='relu'),
    layers.Dense(84, activation='relu'),
    layers.Dense(num_classes, activation='softmax')
])

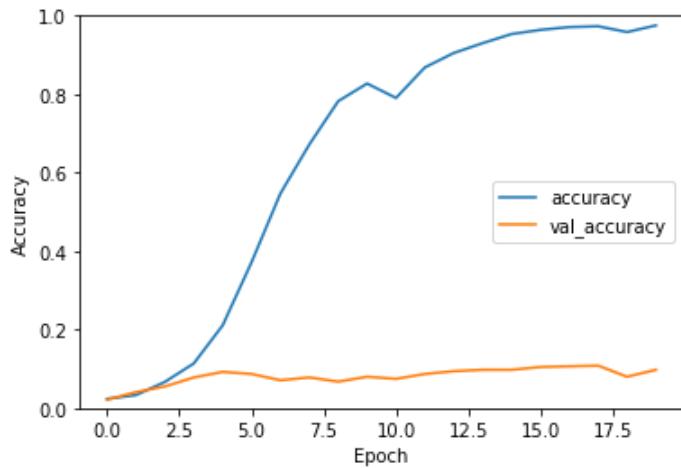
# Compile the model
model.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

# Train the model
history = model.fit(
    X_train, y_train,
    epochs=20,
    batch_size=128,
    verbose=1,
    validation_data=(X_val, y_val)
)

plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0, 1])
plt.legend(loc='center right')
```

```
Epoch 1/20
36/36 [=====] - 61s 2s/step - loss: 4.0446 - accuracy: 0.0235 - val_l
oss: 3.8281 - val_accuracy: 0.0213
Epoch 2/20
36/36 [=====] - 61s 2s/step - loss: 3.8046 - accuracy: 0.0330 - val_l
oss: 3.7937 - val_accuracy: 0.0408
Epoch 3/20
36/36 [=====] - 66s 2s/step - loss: 3.7191 - accuracy: 0.0667 - val_l
oss: 3.7697 - val_accuracy: 0.0550
Epoch 4/20
36/36 [=====] - 69s 2s/step - loss: 3.5514 - accuracy: 0.1137 - val_l
oss: 3.8276 - val_accuracy: 0.0780
Epoch 5/20
36/36 [=====] - 69s 2s/step - loss: 3.1494 - accuracy: 0.2108 - val_l
oss: 3.8722 - val_accuracy: 0.0922
Epoch 6/20
36/36 [=====] - 70s 2s/step - loss: 2.5277 - accuracy: 0.3717 - val_l
oss: 4.2708 - val_accuracy: 0.0869
Epoch 7/20
36/36 [=====] - 70s 2s/step - loss: 1.8446 - accuracy: 0.5470 - val_l
oss: 5.3815 - val_accuracy: 0.0709
Epoch 8/20
36/36 [=====] - 69s 2s/step - loss: 1.3209 - accuracy: 0.6718 - val_l
oss: 6.1650 - val_accuracy: 0.0780
Epoch 9/20
36/36 [=====] - 69s 2s/step - loss: 0.9263 - accuracy: 0.7821 - val_l
oss: 7.1192 - val_accuracy: 0.0674
Epoch 10/20
36/36 [=====] - 74s 2s/step - loss: 0.7416 - accuracy: 0.8271 - val_l
oss: 8.5597 - val_accuracy: 0.0798
Epoch 11/20
36/36 [=====] - 74s 2s/step - loss: 0.9566 - accuracy: 0.7903 - val_l
oss: 8.1215 - val_accuracy: 0.0745
Epoch 12/20
36/36 [=====] - 73s 2s/step - loss: 0.5742 - accuracy: 0.8681 - val_l
oss: 9.8621 - val_accuracy: 0.0869
Epoch 13/20
36/36 [=====] - 72s 2s/step - loss: 0.4384 - accuracy: 0.9049 - val_l
oss: 9.6499 - val_accuracy: 0.0940
Epoch 14/20
36/36 [=====] - 74s 2s/step - loss: 0.3155 - accuracy: 0.9300 - val_l
oss: 10.1538 - val_accuracy: 0.0975
Epoch 15/20
36/36 [=====] - 75s 2s/step - loss: 0.2114 - accuracy: 0.9532 - val_l
oss: 10.5784 - val_accuracy: 0.0975
Epoch 16/20
36/36 [=====] - 70s 2s/step - loss: 0.1713 - accuracy: 0.9639 - val_l
oss: 11.3988 - val_accuracy: 0.1046
Epoch 17/20
36/36 [=====] - 70s 2s/step - loss: 0.1490 - accuracy: 0.9710 - val_l
oss: 11.6427 - val_accuracy: 0.1064
Epoch 18/20
36/36 [=====] - 67s 2s/step - loss: 0.1480 - accuracy: 0.9730 - val_l
oss: 12.1059 - val_accuracy: 0.1082
Epoch 19/20
36/36 [=====] - 66s 2s/step - loss: 0.2020 - accuracy: 0.9583 - val_l
oss: 11.6016 - val_accuracy: 0.0798
Epoch 20/20
36/36 [=====] - 69s 2s/step - loss: 0.1253 - accuracy: 0.9752 - val_l
oss: 12.0452 - val_accuracy: 0.0975
```

Out[10]: <matplotlib.legend.Legend at 0x1db8038e790>



```
In [12]: # Evaluate the model on the testing set
score = model.evaluate(X_test, y_test, verbose=2)
print('Test Loss:', score[0])
print('Test accuracy:', score[1])

18/18 - 2s - loss: 13.3695 - accuracy: 0.0851 - 2s/epoch - 115ms/step
Test Loss: 13.369486808776855
Test accuracy: 0.08510638028383255
```

```
In [15]: # Predict classes for the testing set
y_pred = model.predict(X_test)
y_pred = np.argmax(y_pred, axis=1)

# Calculate precision, recall, and other metrics
print("Classification Report:")
print(classification_report(y_test, y_pred))

# Calculate and plot confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion Matrix')
plt.colorbar()
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.xticks(np.arange(len(class_names)), class_names, rotation=90)
plt.yticks(np.arange(len(class_names)), class_names)
plt.show()
```

18/18 [=====] - 2s 107ms/step

Classification Report:

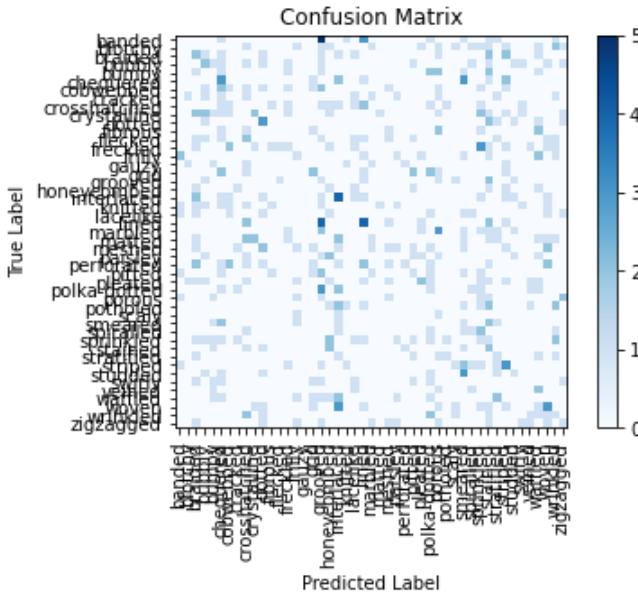
	precision	recall	f1-score	support
0	0.00	0.00	0.00	15
1	0.00	0.00	0.00	16
2	0.10	0.14	0.11	14
3	0.13	0.13	0.13	15
4	0.00	0.00	0.00	12
5	0.11	0.23	0.15	13
6	0.10	0.07	0.08	14
7	0.00	0.00	0.00	10
8	0.00	0.00	0.00	14
9	0.33	0.11	0.16	19
10	0.19	0.30	0.23	10
11	0.00	0.00	0.00	9
12	0.00	0.00	0.00	13
13	0.11	0.06	0.08	17
14	0.11	0.09	0.10	11
15	0.00	0.00	0.00	7
16	0.18	0.17	0.17	12
17	0.03	0.11	0.05	9
18	0.07	0.12	0.09	8
19	0.16	0.25	0.20	16
20	0.00	0.00	0.00	12
21	0.00	0.00	0.00	10
22	0.19	0.27	0.22	15
23	0.00	0.00	0.00	12
24	0.00	0.00	0.00	15
25	0.11	0.06	0.08	16
26	0.00	0.00	0.00	10
27	0.33	0.06	0.10	17
28	0.09	0.09	0.09	11
29	0.29	0.12	0.17	16
30	0.05	0.08	0.06	12
31	0.00	0.00	0.00	9
32	0.00	0.00	0.00	12
33	0.00	0.00	0.00	3
34	0.05	0.11	0.07	9
35	0.14	0.12	0.13	8
36	0.05	0.06	0.05	17
37	0.06	0.22	0.10	9
38	0.08	0.12	0.10	8
39	0.16	0.23	0.19	13
40	0.14	0.11	0.12	9
41	0.00	0.00	0.00	9
42	0.00	0.00	0.00	6
43	0.00	0.00	0.00	12
44	0.19	0.19	0.19	16
45	0.07	0.09	0.08	11
46	0.11	0.08	0.09	13
accuracy			0.09	564
macro avg	0.08	0.08	0.07	564
weighted avg	0.09	0.09	0.08	564

C:\Users\josep\anaconda3\lib\site-packages\sklearn\metrics_classification.py:1245: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))
C:\Users\josep\anaconda3\lib\site-packages\sklearn\metrics_classification.py:1245: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))
C:\Users\josep\anaconda3\lib\site-packages\sklearn\metrics_classification.py:1245: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))



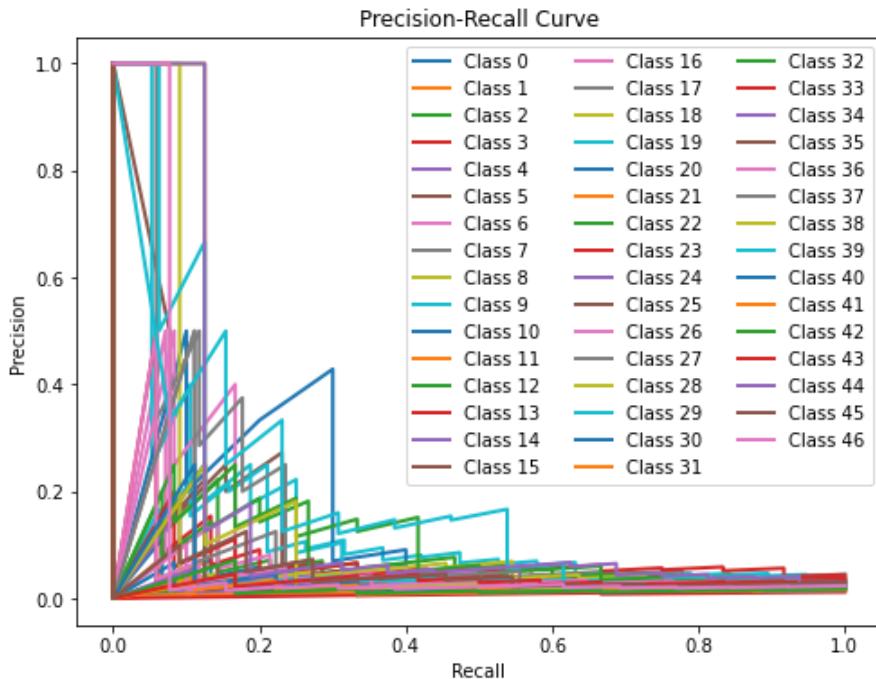
```
In [27]: # Obtain the predicted probabilities for each class
y_pred_prob = model.predict(X_test)

# Binarize the true labels
y_test_bin = label_binarize(y_test, classes=np.unique(y_test))

# Calculate precision and recall for each class
precision = dict()
recall = dict()
for i in range(num_classes):
    precision[i], recall[i], _ = precision_recall_curve(y_test_bin[:, i], y_pred_prob[:, i])

# Plot precision-recall curve for each class
plt.figure(figsize=(8, 6))
for i in range(num_classes):
    plt.plot(recall[i], precision[i], lw=2, label=f'Class {i}')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.legend(loc='upper right', bbox_to_anchor=(1.0, 1.0), ncol=3)
plt.show()
```

18/18 [=====] - 2s 118ms/step



CNN Model

CNN Model with 20 Epochs

```
In [28]: # Define the CNN model architecture
model = keras.Sequential([
    layers.Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(224, 224, 3)),
    layers.MaxPooling2D(pool_size=(2, 2)),
    layers.Conv2D(64, kernel_size=(3, 3), activation='relu'),
    layers.MaxPooling2D(pool_size=(2, 2)),
    layers.Conv2D(128, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(256, activation='relu', kernel_regularizer=regularizers.l2(0.001)),
    layers.Dropout(0.5),
    layers.Dense(128, activation='relu', kernel_regularizer=regularizers.l2(0.001)),
    layers.Dropout(0.5),
    layers.Dense(num_classes, activation='softmax')
])

# Compile the model
model.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

# Train the model
history = model.fit(
    X_train, y_train,
    epochs=20,
    batch_size=128,
    verbose=1,
    validation_data=(X_val, y_val)
)
```

```
Epoch 1/20
36/36 [=====] - 156s 4s/step - loss: 4.4337 - accuracy: 0.0224 - val_
loss: 4.1773 - val_accuracy: 0.0408
Epoch 2/20
36/36 [=====] - 164s 5s/step - loss: 4.0893 - accuracy: 0.0268 - val_
loss: 4.0152 - val_accuracy: 0.0266
Epoch 3/20
36/36 [=====] - 164s 5s/step - loss: 3.9661 - accuracy: 0.0299 - val_
loss: 3.9531 - val_accuracy: 0.0301
Epoch 4/20
36/36 [=====] - 165s 5s/step - loss: 3.9176 - accuracy: 0.0370 - val_
loss: 3.9213 - val_accuracy: 0.0319
Epoch 5/20
36/36 [=====] - 157s 4s/step - loss: 3.8951 - accuracy: 0.0397 - val_
loss: 3.9102 - val_accuracy: 0.0337
Epoch 6/20
36/36 [=====] - 152s 4s/step - loss: 3.8761 - accuracy: 0.0401 - val_
loss: 3.9066 - val_accuracy: 0.0461
Epoch 7/20
36/36 [=====] - 155s 4s/step - loss: 3.8623 - accuracy: 0.0468 - val_
loss: 3.8946 - val_accuracy: 0.0514
Epoch 8/20
36/36 [=====] - 157s 4s/step - loss: 3.8446 - accuracy: 0.0543 - val_
loss: 3.8947 - val_accuracy: 0.0390
Epoch 9/20
36/36 [=====] - 159s 4s/step - loss: 3.8133 - accuracy: 0.0516 - val_
loss: 3.8855 - val_accuracy: 0.0496
Epoch 10/20
36/36 [=====] - 163s 5s/step - loss: 3.7857 - accuracy: 0.0678 - val_
loss: 3.8861 - val_accuracy: 0.0550
Epoch 11/20
36/36 [=====] - 171s 5s/step - loss: 3.7723 - accuracy: 0.0731 - val_
loss: 3.8997 - val_accuracy: 0.0603
Epoch 12/20
36/36 [=====] - 170s 5s/step - loss: 3.7425 - accuracy: 0.0889 - val_
loss: 3.9206 - val_accuracy: 0.0638
Epoch 13/20
36/36 [=====] - 167s 5s/step - loss: 3.7133 - accuracy: 0.1053 - val_
loss: 3.9079 - val_accuracy: 0.0479
Epoch 14/20
36/36 [=====] - 168s 5s/step - loss: 3.6562 - accuracy: 0.1197 - val_
loss: 3.9415 - val_accuracy: 0.0745
Epoch 15/20
36/36 [=====] - 163s 5s/step - loss: 3.5938 - accuracy: 0.1463 - val_
loss: 3.9703 - val_accuracy: 0.0833
Epoch 16/20
36/36 [=====] - 159s 4s/step - loss: 3.5603 - accuracy: 0.1653 - val_
loss: 3.9865 - val_accuracy: 0.0851
Epoch 17/20
36/36 [=====] - 157s 4s/step - loss: 3.4730 - accuracy: 0.1886 - val_
loss: 4.0279 - val_accuracy: 0.0851
Epoch 18/20
36/36 [=====] - 158s 4s/step - loss: 3.3745 - accuracy: 0.2152 - val_
loss: 4.0710 - val_accuracy: 0.0869
Epoch 19/20
36/36 [=====] - 159s 4s/step - loss: 3.3252 - accuracy: 0.2431 - val_
loss: 4.0789 - val_accuracy: 0.0869
Epoch 20/20
36/36 [=====] - 150s 4s/step - loss: 3.2689 - accuracy: 0.2538 - val_
loss: 4.0851 - val_accuracy: 0.1082
```

CNN Model with 50 Epochs

```
In [36]: # Now with 50 epochs
model = keras.Sequential([
    layers.Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(224, 224, 3)),
    layers.MaxPooling2D(pool_size=(2, 2)),
    layers.Conv2D(64, kernel_size=(3, 3), activation='relu'),
    layers.MaxPooling2D(pool_size=(2, 2)),
    layers.Conv2D(128, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(256, activation='relu', kernel_regularizer=regularizers.l2(0.001)),
    layers.Dropout(0.5),
    layers.Dense(128, activation='relu', kernel_regularizer=regularizers.l2(0.001)),
    layers.Dropout(0.5),
    layers.Dense(num_classes, activation='softmax')
])

# Compile the model
model.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

# Train the model
history = model.fit(
    X_train, y_train,
    epochs=50,
    batch_size=128,
    verbose=1,
    validation_data=(X_val, y_val)
)

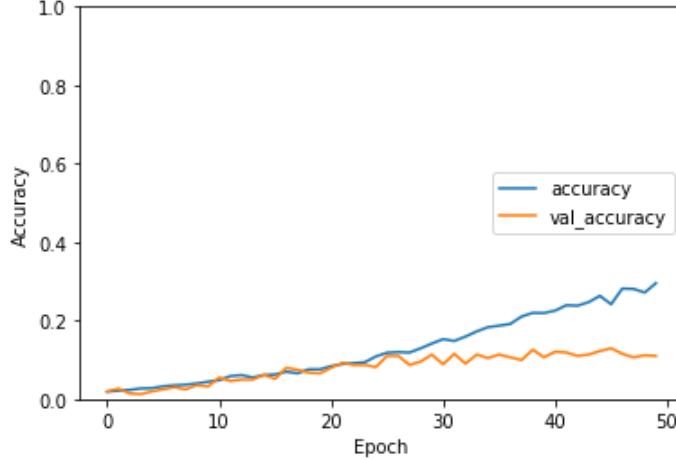
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0, 1])
plt.legend(loc='center right')
```

```
Epoch 1/50
36/36 [=====] - 149s 4s/step - loss: 4.4766 - accuracy: 0.0188 - val_
loss: 4.1788 - val_accuracy: 0.0195
Epoch 2/50
36/36 [=====] - 156s 4s/step - loss: 4.0969 - accuracy: 0.0215 - val_
loss: 4.0316 - val_accuracy: 0.0266
Epoch 3/50
36/36 [=====] - 156s 4s/step - loss: 3.9966 - accuracy: 0.0233 - val_
loss: 3.9690 - val_accuracy: 0.0142
Epoch 4/50
36/36 [=====] - 155s 4s/step - loss: 3.9480 - accuracy: 0.0270 - val_
loss: 3.9328 - val_accuracy: 0.0124
Epoch 5/50
36/36 [=====] - 156s 4s/step - loss: 3.9111 - accuracy: 0.0277 - val_
loss: 3.8988 - val_accuracy: 0.0195
Epoch 6/50
36/36 [=====] - 157s 4s/step - loss: 3.8857 - accuracy: 0.0326 - val_
loss: 3.8806 - val_accuracy: 0.0248
Epoch 7/50
36/36 [=====] - 156s 4s/step - loss: 3.8592 - accuracy: 0.0350 - val_
loss: 3.8789 - val_accuracy: 0.0301
Epoch 8/50
36/36 [=====] - 155s 4s/step - loss: 3.8459 - accuracy: 0.0366 - val_
loss: 3.8480 - val_accuracy: 0.0248
Epoch 9/50
36/36 [=====] - 159s 4s/step - loss: 3.8149 - accuracy: 0.0399 - val_
loss: 3.8382 - val_accuracy: 0.0355
Epoch 10/50
36/36 [=====] - 161s 4s/step - loss: 3.7899 - accuracy: 0.0441 - val_
loss: 3.8075 - val_accuracy: 0.0319
Epoch 11/50
36/36 [=====] - 153s 4s/step - loss: 3.7721 - accuracy: 0.0481 - val_
loss: 3.7778 - val_accuracy: 0.0550
Epoch 12/50
36/36 [=====] - 175s 5s/step - loss: 3.7468 - accuracy: 0.0585 - val_
loss: 3.7511 - val_accuracy: 0.0461
Epoch 13/50
36/36 [=====] - 157s 4s/step - loss: 3.7195 - accuracy: 0.0612 - val_
loss: 3.7508 - val_accuracy: 0.0496
Epoch 14/50
36/36 [=====] - 166s 5s/step - loss: 3.6933 - accuracy: 0.0550 - val_
loss: 3.7239 - val_accuracy: 0.0496
Epoch 15/50
36/36 [=====] - 173s 5s/step - loss: 3.6571 - accuracy: 0.0614 - val_
loss: 3.7082 - val_accuracy: 0.0621
Epoch 16/50
36/36 [=====] - 165s 5s/step - loss: 3.6413 - accuracy: 0.0627 - val_
loss: 3.6513 - val_accuracy: 0.0514
Epoch 17/50
36/36 [=====] - 166s 5s/step - loss: 3.6190 - accuracy: 0.0698 - val_
loss: 3.6453 - val_accuracy: 0.0798
Epoch 18/50
36/36 [=====] - 167s 5s/step - loss: 3.6045 - accuracy: 0.0656 - val_
loss: 3.6323 - val_accuracy: 0.0745
Epoch 19/50
36/36 [=====] - 169s 5s/step - loss: 3.5804 - accuracy: 0.0758 - val_
loss: 3.7041 - val_accuracy: 0.0674
Epoch 20/50
36/36 [=====] - 155s 4s/step - loss: 3.5598 - accuracy: 0.0758 - val_
loss: 3.6381 - val_accuracy: 0.0656
Epoch 21/50
36/36 [=====] - 157s 4s/step - loss: 3.5369 - accuracy: 0.0840 - val_
loss: 3.6220 - val_accuracy: 0.0798
Epoch 22/50
36/36 [=====] - 157s 4s/step - loss: 3.4940 - accuracy: 0.0891 - val_
loss: 3.6048 - val_accuracy: 0.0922
Epoch 23/50
36/36 [=====] - 147s 4s/step - loss: 3.4633 - accuracy: 0.0915 - val_
```

```
loss: 3.6345 - val_accuracy: 0.0869
Epoch 24/50
36/36 [=====] - 152s 4s/step - loss: 3.4180 - accuracy: 0.0938 - val_
loss: 3.6802 - val_accuracy: 0.0869
Epoch 25/50
36/36 [=====] - 156s 4s/step - loss: 3.3909 - accuracy: 0.1086 - val_
loss: 3.6230 - val_accuracy: 0.0816
Epoch 26/50
36/36 [=====] - 156s 4s/step - loss: 3.3713 - accuracy: 0.1181 - val_
loss: 3.6291 - val_accuracy: 0.1099
Epoch 27/50
36/36 [=====] - 156s 4s/step - loss: 3.3405 - accuracy: 0.1197 - val_
loss: 3.6549 - val_accuracy: 0.1099
Epoch 28/50
36/36 [=====] - 157s 4s/step - loss: 3.3247 - accuracy: 0.1186 - val_
loss: 3.6438 - val_accuracy: 0.0869
Epoch 29/50
36/36 [=====] - 156s 4s/step - loss: 3.2885 - accuracy: 0.1294 - val_
loss: 3.6316 - val_accuracy: 0.0957
Epoch 30/50
36/36 [=====] - 156s 4s/step - loss: 3.2517 - accuracy: 0.1416 - val_
loss: 3.6408 - val_accuracy: 0.1135
Epoch 31/50
36/36 [=====] - 156s 4s/step - loss: 3.1863 - accuracy: 0.1527 - val_
loss: 3.7325 - val_accuracy: 0.0887
Epoch 32/50
36/36 [=====] - 156s 4s/step - loss: 3.1833 - accuracy: 0.1480 - val_
loss: 3.6641 - val_accuracy: 0.1152
Epoch 33/50
36/36 [=====] - 155s 4s/step - loss: 3.1235 - accuracy: 0.1594 - val_
loss: 3.6873 - val_accuracy: 0.0904
Epoch 34/50
36/36 [=====] - 151s 4s/step - loss: 3.1272 - accuracy: 0.1724 - val_
loss: 3.7327 - val_accuracy: 0.1135
Epoch 35/50
36/36 [=====] - 148s 4s/step - loss: 3.0488 - accuracy: 0.1833 - val_
loss: 3.7526 - val_accuracy: 0.1046
Epoch 36/50
36/36 [=====] - 155s 4s/step - loss: 2.9972 - accuracy: 0.1871 - val_
loss: 3.8094 - val_accuracy: 0.1135
Epoch 37/50
36/36 [=====] - 160s 4s/step - loss: 3.0107 - accuracy: 0.1915 - val_
loss: 3.6902 - val_accuracy: 0.1064
Epoch 38/50
36/36 [=====] - 161s 4s/step - loss: 2.9664 - accuracy: 0.2103 - val_
loss: 3.7438 - val_accuracy: 0.0993
Epoch 39/50
36/36 [=====] - 156s 4s/step - loss: 2.9222 - accuracy: 0.2196 - val_
loss: 3.7288 - val_accuracy: 0.1259
Epoch 40/50
36/36 [=====] - 158s 4s/step - loss: 2.9081 - accuracy: 0.2192 - val_
loss: 3.6928 - val_accuracy: 0.1064
Epoch 41/50
36/36 [=====] - 156s 4s/step - loss: 2.8655 - accuracy: 0.2254 - val_
loss: 3.8066 - val_accuracy: 0.1206
Epoch 42/50
36/36 [=====] - 155s 4s/step - loss: 2.8092 - accuracy: 0.2391 - val_
loss: 3.8954 - val_accuracy: 0.1188
Epoch 43/50
36/36 [=====] - 156s 4s/step - loss: 2.8246 - accuracy: 0.2380 - val_
loss: 3.8362 - val_accuracy: 0.1099
Epoch 44/50
36/36 [=====] - 156s 4s/step - loss: 2.7875 - accuracy: 0.2471 - val_
loss: 3.8609 - val_accuracy: 0.1135
Epoch 45/50
36/36 [=====] - 149s 4s/step - loss: 2.7892 - accuracy: 0.2631 - val_
loss: 3.8135 - val_accuracy: 0.1223
Epoch 46/50
```

```
36/36 [=====] - 145s 4s/step - loss: 2.7752 - accuracy: 0.2416 - val_loss: 3.9031 - val_accuracy: 0.1294
Epoch 47/50
36/36 [=====] - 148s 4s/step - loss: 2.6770 - accuracy: 0.2817 - val_loss: 3.9568 - val_accuracy: 0.1152
Epoch 48/50
36/36 [=====] - 156s 4s/step - loss: 2.6802 - accuracy: 0.2808 - val_loss: 3.9217 - val_accuracy: 0.1064
Epoch 49/50
36/36 [=====] - 160s 4s/step - loss: 2.7268 - accuracy: 0.2715 - val_loss: 4.0463 - val_accuracy: 0.1117
Epoch 50/50
36/36 [=====] - 156s 4s/step - loss: 2.6147 - accuracy: 0.2954 - val_loss: 3.9666 - val_accuracy: 0.1099
```

Out[36]: <matplotlib.legend.Legend at 0x1dd505dc820>



```
In [78]: # Evaluate the model on the testing set
score = model.evaluate(X_test, y_test, verbose=2)
print('Test Loss:', score[0])
print('Test accuracy:', score[1])
```

```
18/18 - 6s - loss: 3.5595 - accuracy: 0.0780 - 6s/epoch - 318ms/step
Test Loss: 3.5595102310180664
Test accuracy: 0.07801418751478195
```

AlexNet Model

AlexNet Model with 50 Epochs and Learning Rate 0.001

```
In [34]: model = keras.models.Sequential([
    keras.layers.Conv2D(filters=96, kernel_size=(11,11), strides=(4,4), activation='relu', input_shape=(224,224,3)),
    keras.layers.BatchNormalization(),
    keras.layers.MaxPool2D(pool_size=(3,3), strides=(2,2)),
    keras.layers.Conv2D(filters=256, kernel_size=(5,5), strides=(1,1), activation='relu', padding='same'),
    keras.layers.BatchNormalization(),
    keras.layers.MaxPool2D(pool_size=(3,3), strides=(2,2)),
    keras.layers.Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), activation='relu', padding='same'),
    keras.layers.BatchNormalization(),
    keras.layers.Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), activation='relu', padding='same'),
    keras.layers.BatchNormalization(),
    keras.layers.Conv2D(filters=256, kernel_size=(3,3), strides=(1,1), activation='relu', padding='same'),
    keras.layers.BatchNormalization(),
    keras.layers.MaxPool2D(pool_size=(3,3), strides=(2,2)),
    keras.layers.Flatten(),
    keras.layers.Dense(4096, activation='relu'),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(4096, activation='relu'),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(num_classes, activation='softmax')
])

model.compile(loss='categorical_crossentropy', optimizer=tf.optimizers.SGD(learning_rate=0.001))
model.summary()

# Train the model
history = model.fit(
    X_train, y_train,
    epochs=50,
    verbose=1,
    validation_data=(X_val, y_val)
)

plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0, 1])
plt.legend(loc='lower right')
plt.show()
```

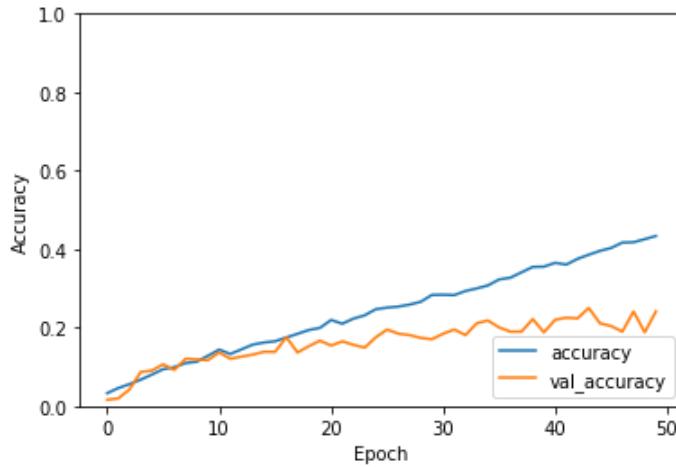
Model: "sequential_2"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_10 (Conv2D)	(None, 54, 54, 96)	34944
batch_normalization_10 (BatchNormalization)	(None, 54, 54, 96)	384
max_pooling2d_6 (MaxPooling2D)	(None, 26, 26, 96)	0
conv2d_11 (Conv2D)	(None, 26, 26, 256)	614656
batch_normalization_11 (BatchNormalization)	(None, 26, 26, 256)	1024
max_pooling2d_7 (MaxPooling2D)	(None, 12, 12, 256)	0
conv2d_12 (Conv2D)	(None, 12, 12, 384)	885120
batch_normalization_12 (BatchNormalization)	(None, 12, 12, 384)	1536
conv2d_13 (Conv2D)	(None, 12, 12, 384)	1327488
batch_normalization_13 (BatchNormalization)	(None, 12, 12, 384)	1536
conv2d_14 (Conv2D)	(None, 12, 12, 256)	884992
batch_normalization_14 (BatchNormalization)	(None, 12, 12, 256)	1024
max_pooling2d_8 (MaxPooling2D)	(None, 5, 5, 256)	0
flatten_2 (Flatten)	(None, 6400)	0
dense_6 (Dense)	(None, 4096)	26218496
dropout_4 (Dropout)	(None, 4096)	0
dense_7 (Dense)	(None, 4096)	16781312
dropout_5 (Dropout)	(None, 4096)	0
dense_8 (Dense)	(None, 47)	192559
<hr/>		
Total params: 46,945,071		
Trainable params: 46,942,319		
Non-trainable params: 2,752		
<hr/>		
Epoch 1/50		
141/141	[=====]	- 214s 2s/step - loss: 5.2604 - accuracy: 0.0328 - val_loss: 3.9002 - val_accuracy: 0.0160
Epoch 2/50		
141/141	[=====]	- 219s 2s/step - loss: 4.5072 - accuracy: 0.0459 - val_loss: 3.8646 - val_accuracy: 0.0195
Epoch 3/50		
141/141	[=====]	- 218s 2s/step - loss: 4.1889 - accuracy: 0.0559 - val_loss: 3.7730 - val_accuracy: 0.0426
Epoch 4/50		
141/141	[=====]	- 218s 2s/step - loss: 4.0252 - accuracy: 0.0674 - val_loss: 3.6270 - val_accuracy: 0.0869
Epoch 5/50		

```
141/141 [=====] - 220s 2s/step - loss: 3.8617 - accuracy: 0.0802 - va  
l_loss: 3.6061 - val_accuracy: 0.0904  
Epoch 6/50  
141/141 [=====] - 219s 2s/step - loss: 3.7654 - accuracy: 0.0940 - va  
l_loss: 3.5801 - val_accuracy: 0.1064  
Epoch 7/50  
141/141 [=====] - 218s 2s/step - loss: 3.6971 - accuracy: 0.0991 - va  
l_loss: 3.6370 - val_accuracy: 0.0922  
Epoch 8/50  
141/141 [=====] - 211s 1s/step - loss: 3.6016 - accuracy: 0.1095 - va  
l_loss: 3.5480 - val_accuracy: 0.1206  
Epoch 9/50  
141/141 [=====] - 204s 1s/step - loss: 3.5590 - accuracy: 0.1130 - va  
l_loss: 3.4531 - val_accuracy: 0.1188  
Epoch 10/50  
141/141 [=====] - 222s 2s/step - loss: 3.4700 - accuracy: 0.1281 - va  
l_loss: 3.5021 - val_accuracy: 0.1170  
Epoch 11/50  
141/141 [=====] - 223s 2s/step - loss: 3.4365 - accuracy: 0.1438 - va  
l_loss: 3.4148 - val_accuracy: 0.1365  
Epoch 12/50  
141/141 [=====] - 219s 2s/step - loss: 3.4060 - accuracy: 0.1328 - va  
l_loss: 3.4566 - val_accuracy: 0.1206  
Epoch 13/50  
141/141 [=====] - 219s 2s/step - loss: 3.3826 - accuracy: 0.1445 - va  
l_loss: 3.4342 - val_accuracy: 0.1259  
Epoch 14/50  
141/141 [=====] - 217s 2s/step - loss: 3.3205 - accuracy: 0.1567 - va  
l_loss: 3.4636 - val_accuracy: 0.1312  
Epoch 15/50  
141/141 [=====] - 220s 2s/step - loss: 3.2457 - accuracy: 0.1620 - va  
l_loss: 3.4046 - val_accuracy: 0.1383  
Epoch 16/50  
141/141 [=====] - 222s 2s/step - loss: 3.2403 - accuracy: 0.1651 - va  
l_loss: 3.4926 - val_accuracy: 0.1383  
Epoch 17/50  
141/141 [=====] - 207s 1s/step - loss: 3.2008 - accuracy: 0.1749 - va  
l_loss: 3.2705 - val_accuracy: 0.1755  
Epoch 18/50  
141/141 [=====] - 214s 2s/step - loss: 3.1591 - accuracy: 0.1842 - va  
l_loss: 3.4299 - val_accuracy: 0.1365  
Epoch 19/50  
141/141 [=====] - 218s 2s/step - loss: 3.1296 - accuracy: 0.1937 - va  
l_loss: 3.3244 - val_accuracy: 0.1525  
Epoch 20/50  
141/141 [=====] - 218s 2s/step - loss: 3.0948 - accuracy: 0.1992 - va  
l_loss: 3.3670 - val_accuracy: 0.1667  
Epoch 21/50  
141/141 [=====] - 220s 2s/step - loss: 3.0409 - accuracy: 0.2194 - va  
l_loss: 3.3938 - val_accuracy: 0.1543  
Epoch 22/50  
141/141 [=====] - 218s 2s/step - loss: 3.0080 - accuracy: 0.2099 - va  
l_loss: 3.2298 - val_accuracy: 0.1649  
Epoch 23/50  
141/141 [=====] - 218s 2s/step - loss: 2.9915 - accuracy: 0.2230 - va  
l_loss: 3.3630 - val_accuracy: 0.1560  
Epoch 24/50  
141/141 [=====] - 217s 2s/step - loss: 2.9364 - accuracy: 0.2314 - va  
l_loss: 3.5324 - val_accuracy: 0.1489  
Epoch 25/50  
141/141 [=====] - 204s 1s/step - loss: 2.9078 - accuracy: 0.2469 - va  
l_loss: 3.2160 - val_accuracy: 0.1755  
Epoch 26/50  
141/141 [=====] - 209s 1s/step - loss: 2.8707 - accuracy: 0.2509 - va  
l_loss: 3.2589 - val_accuracy: 0.1950  
Epoch 27/50  
141/141 [=====] - 217s 2s/step - loss: 2.8422 - accuracy: 0.2535 - va  
l_loss: 3.1965 - val_accuracy: 0.1844
```

```
Epoch 28/50
141/141 [=====] - 219s 2s/step - loss: 2.8176 - accuracy: 0.2586 - val_loss: 3.2422 - val_accuracy: 0.1809
Epoch 29/50
141/141 [=====] - 217s 2s/step - loss: 2.7699 - accuracy: 0.2657 - val_loss: 3.2973 - val_accuracy: 0.1738
Epoch 30/50
141/141 [=====] - 218s 2s/step - loss: 2.7338 - accuracy: 0.2832 - val_loss: 3.2957 - val_accuracy: 0.1702
Epoch 31/50
141/141 [=====] - 220s 2s/step - loss: 2.6969 - accuracy: 0.2837 - val_loss: 3.3316 - val_accuracy: 0.1844
Epoch 32/50
141/141 [=====] - 220s 2s/step - loss: 2.6938 - accuracy: 0.2828 - val_loss: 3.1074 - val_accuracy: 0.1950
Epoch 33/50
141/141 [=====] - 210s 1s/step - loss: 2.6624 - accuracy: 0.2932 - val_loss: 3.2313 - val_accuracy: 0.1809
Epoch 34/50
141/141 [=====] - 204s 1s/step - loss: 2.6107 - accuracy: 0.2999 - val_loss: 3.0846 - val_accuracy: 0.2110
Epoch 35/50
141/141 [=====] - 218s 2s/step - loss: 2.5728 - accuracy: 0.3074 - val_loss: 3.1038 - val_accuracy: 0.2181
Epoch 36/50
141/141 [=====] - 219s 2s/step - loss: 2.5344 - accuracy: 0.3227 - val_loss: 3.0639 - val_accuracy: 0.2004
Epoch 37/50
141/141 [=====] - 219s 2s/step - loss: 2.5124 - accuracy: 0.3271 - val_loss: 3.2038 - val_accuracy: 0.1897
Epoch 38/50
141/141 [=====] - 219s 2s/step - loss: 2.4633 - accuracy: 0.3409 - val_loss: 3.1867 - val_accuracy: 0.1897
Epoch 39/50
141/141 [=====] - 218s 2s/step - loss: 2.4330 - accuracy: 0.3546 - val_loss: 3.0909 - val_accuracy: 0.2216
Epoch 40/50
141/141 [=====] - 217s 2s/step - loss: 2.3823 - accuracy: 0.3553 - val_loss: 3.2403 - val_accuracy: 0.1879
Epoch 41/50
141/141 [=====] - 217s 2s/step - loss: 2.3737 - accuracy: 0.3650 - val_loss: 3.0918 - val_accuracy: 0.2199
Epoch 42/50
141/141 [=====] - 203s 1s/step - loss: 2.3417 - accuracy: 0.3606 - val_loss: 3.0464 - val_accuracy: 0.2252
Epoch 43/50
141/141 [=====] - 212s 2s/step - loss: 2.2965 - accuracy: 0.3752 - val_loss: 3.0277 - val_accuracy: 0.2234
Epoch 44/50
141/141 [=====] - 218s 2s/step - loss: 2.2651 - accuracy: 0.3856 - val_loss: 2.9942 - val_accuracy: 0.2500
Epoch 45/50
141/141 [=====] - 217s 2s/step - loss: 2.2616 - accuracy: 0.3956 - val_loss: 3.1048 - val_accuracy: 0.2110
Epoch 46/50
141/141 [=====] - 218s 2s/step - loss: 2.1995 - accuracy: 0.4029 - val_loss: 3.0981 - val_accuracy: 0.2039
Epoch 47/50
141/141 [=====] - 217s 2s/step - loss: 2.1449 - accuracy: 0.4167 - val_loss: 3.0866 - val_accuracy: 0.1897
Epoch 48/50
141/141 [=====] - 217s 2s/step - loss: 2.1445 - accuracy: 0.4176 - val_loss: 2.9823 - val_accuracy: 0.2411
Epoch 49/50
141/141 [=====] - 218s 2s/step - loss: 2.0908 - accuracy: 0.4253 - val_loss: 3.5063 - val_accuracy: 0.1879
Epoch 50/50
```

```
141/141 [=====] - 203s 1s/step - loss: 2.0701 - accuracy: 0.4335 - val_loss: 3.0149 - val_accuracy: 0.2411
```



```
In [35]: # Evaluate the model on the testing set
score = model.evaluate(X_test, y_test, verbose=2)
print('Test Loss:', score[0])
print('Test accuracy:', score[1])
```

```
18/18 - 6s - loss: 3.0014 - accuracy: 0.2518 - 6s/epoch - 331ms/step
Test Loss: 3.0014026165008545
Test accuracy: 0.25177305936813354
```

```
In [36]: # Predict classes for the testing set
y_pred = model.predict(X_test)
y_pred = np.argmax(y_pred, axis=1)

# Convert probabilities to class labels
y_test = np.argmax(y_test, axis=1)

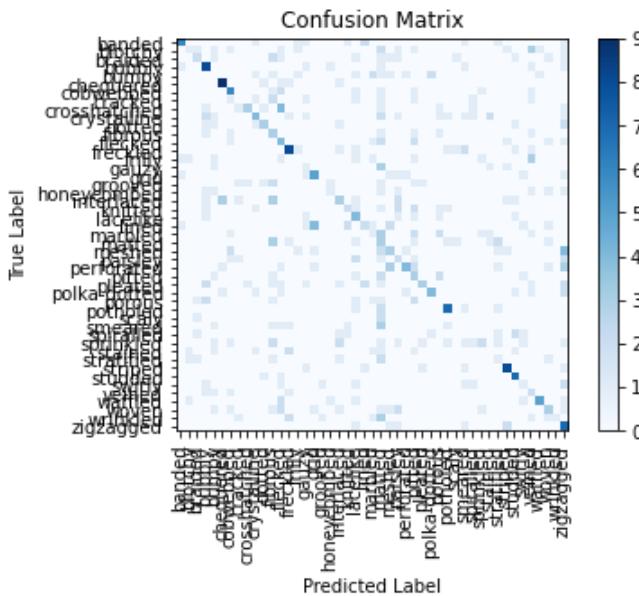
# Calculate precision, recall, and other metrics
print("Classification Report:")
print(classification_report(y_test, y_pred))

# Calculate and plot confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion Matrix')
plt.colorbar()
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.xticks(np.arange(len(class_names)), class_names, rotation=90)
plt.yticks(np.arange(len(class_names)), class_names)
plt.show()
```

18/18 [=====] - 6s 323ms/step

Classification Report:

	precision	recall	f1-score	support
0	0.60	0.40	0.48	15
1	0.20	0.06	0.10	16
2	0.29	0.14	0.19	14
3	0.35	0.53	0.42	15
4	0.00	0.00	0.00	12
5	0.64	0.69	0.67	13
6	0.35	0.43	0.39	14
7	0.00	0.00	0.00	10
8	0.60	0.21	0.32	14
9	0.36	0.21	0.27	19
10	0.43	0.30	0.35	10
11	0.12	0.33	0.18	9
12	0.17	0.31	0.22	13
13	0.47	0.47	0.47	17
14	0.00	0.00	0.00	11
15	0.10	0.14	0.12	7
16	0.42	0.42	0.42	12
17	0.00	0.00	0.00	9
18	0.25	0.12	0.17	8
19	0.25	0.19	0.21	16
20	0.29	0.17	0.21	12
21	0.19	0.40	0.26	10
22	0.29	0.13	0.18	15
23	0.11	0.08	0.10	12
24	0.06	0.20	0.10	15
25	0.23	0.19	0.21	16
26	0.12	0.20	0.15	10
27	0.50	0.24	0.32	17
28	0.13	0.27	0.18	11
29	0.50	0.19	0.27	16
30	0.57	0.33	0.42	12
31	0.00	0.00	0.00	9
32	0.54	0.58	0.56	12
33	0.00	0.00	0.00	3
34	0.00	0.00	0.00	9
35	0.00	0.00	0.00	8
36	0.33	0.12	0.17	17
37	0.00	0.00	0.00	9
38	0.11	0.12	0.12	8
39	0.53	0.62	0.57	13
40	0.44	0.78	0.56	9
41	0.14	0.11	0.12	9
42	0.11	0.33	0.17	6
43	0.71	0.42	0.53	12
44	0.27	0.19	0.22	16
45	1.00	0.09	0.17	11
46	0.19	0.54	0.28	13
accuracy			0.25	564
macro avg	0.28	0.24	0.23	564
weighted avg	0.30	0.25	0.25	564



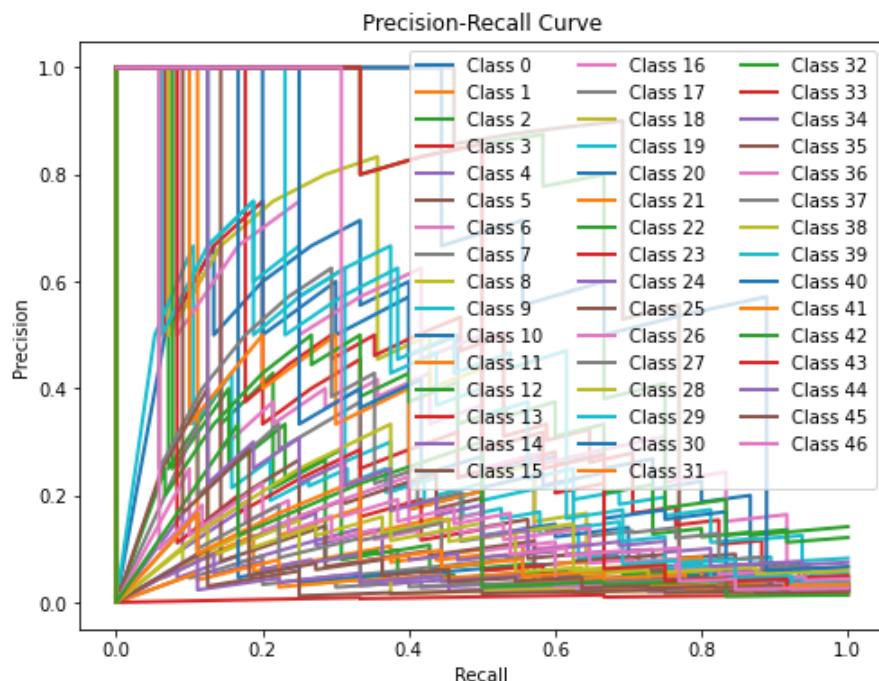
```
In [37]: # Obtain the predicted probabilities for each class
y_pred_prob = model.predict(X_test)

# Binarize the true labels
y_test_bin = label_binarize(y_test, classes=np.unique(y_test))

# Calculate precision and recall for each class
precision = dict()
recall = dict()
for i in range(num_classes):
    precision[i], recall[i], _ = precision_recall_curve(y_test_bin[:, i], y_pred_prob[:, i])

# Plot precision-recall curve for each class
plt.figure(figsize=(8, 6))
for i in range(num_classes):
    plt.plot(recall[i], precision[i], lw=2, label=f'Class {i}')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.legend(loc='upper right', bbox_to_anchor=(1.0, 1.0), ncol=3)
plt.show()
```

18/18 [=====] - 6s 329ms/step



AlexNet Model with 50 Epochs and Learning Rate 0.01

```
In [47]: model = keras.models.Sequential([
    keras.layers.Conv2D(filters=96, kernel_size=(11,11), strides=(4,4), activation='relu', input_shape=(224,224,3)),
    keras.layers.BatchNormalization(),
    keras.layers.MaxPool2D(pool_size=(3,3), strides=(2,2)),
    keras.layers.Conv2D(filters=256, kernel_size=(5,5), strides=(1,1), activation='relu', padding='same'),
    keras.layers.BatchNormalization(),
    keras.layers.MaxPool2D(pool_size=(3,3), strides=(2,2)),
    keras.layers.Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), activation='relu', padding='same'),
    keras.layers.BatchNormalization(),
    keras.layers.Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), activation='relu', padding='same'),
    keras.layers.BatchNormalization(),
    keras.layers.Conv2D(filters=256, kernel_size=(3,3), strides=(1,1), activation='relu', padding='same'),
    keras.layers.BatchNormalization(),
    keras.layers.MaxPool2D(pool_size=(3,3), strides=(2,2)),
    keras.layers.Flatten(),
    keras.layers.Dense(4096, activation='relu'),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(4096, activation='relu'),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(num_classes, activation='softmax')
])

model.compile(loss='categorical_crossentropy', optimizer=tf.optimizers.SGD(learning_rate=0.01))
model.summary()

# Train the model
history = model.fit(
    X_train, y_train,
    epochs=50,
    verbose=1,
    validation_data=(X_val, y_val)
)

plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0, 1])
plt.legend(loc='lower right')
plt.show()
```

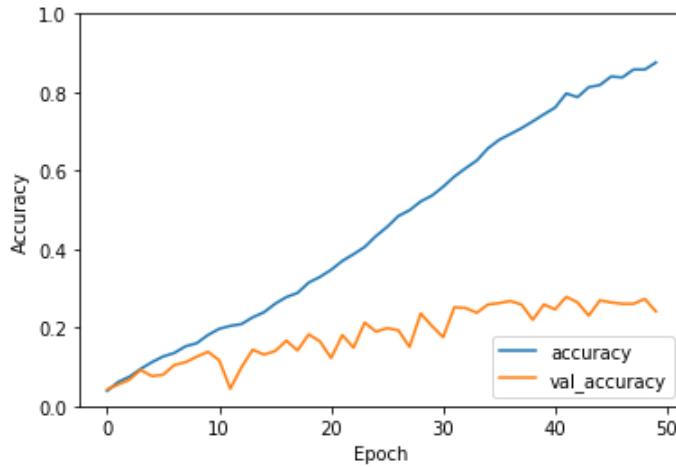
Model: "sequential_3"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_15 (Conv2D)	(None, 54, 54, 96)	34944
batch_normalization_15 (BatchNormalization)	(None, 54, 54, 96)	384
max_pooling2d_9 (MaxPooling2D)	(None, 26, 26, 96)	0
conv2d_16 (Conv2D)	(None, 26, 26, 256)	614656
batch_normalization_16 (BatchNormalization)	(None, 26, 26, 256)	1024
max_pooling2d_10 (MaxPooling2D)	(None, 12, 12, 256)	0
conv2d_17 (Conv2D)	(None, 12, 12, 384)	885120
batch_normalization_17 (BatchNormalization)	(None, 12, 12, 384)	1536
conv2d_18 (Conv2D)	(None, 12, 12, 384)	1327488
batch_normalization_18 (BatchNormalization)	(None, 12, 12, 384)	1536
conv2d_19 (Conv2D)	(None, 12, 12, 256)	884992
batch_normalization_19 (BatchNormalization)	(None, 12, 12, 256)	1024
max_pooling2d_11 (MaxPooling2D)	(None, 5, 5, 256)	0
flatten_3 (Flatten)	(None, 6400)	0
dense_9 (Dense)	(None, 4096)	26218496
dropout_6 (Dropout)	(None, 4096)	0
dense_10 (Dense)	(None, 4096)	16781312
dropout_7 (Dropout)	(None, 4096)	0
dense_11 (Dense)	(None, 47)	192559
<hr/>		
Total params: 46,945,071		
Trainable params: 46,942,319		
Non-trainable params: 2,752		
<hr/>		
Epoch 1/50		
141/141	[=====]	- 215s 2s/step - loss: 4.8836 - accuracy: 0.0388 - val_loss: 3.8321 - val_accuracy: 0.0426
Epoch 2/50		
141/141	[=====]	- 220s 2s/step - loss: 4.0172 - accuracy: 0.0616 - val_loss: 3.8202 - val_accuracy: 0.0550
Epoch 3/50		
141/141	[=====]	- 219s 2s/step - loss: 3.8124 - accuracy: 0.0749 - val_loss: 3.7001 - val_accuracy: 0.0674
Epoch 4/50		
141/141	[=====]	- 218s 2s/step - loss: 3.6461 - accuracy: 0.0946 - val_loss: 3.5628 - val_accuracy: 0.0922
Epoch 5/50		

```
141/141 [=====] - 220s 2s/step - loss: 3.5394 - accuracy: 0.1119 - va  
l_loss: 3.6397 - val_accuracy: 0.0762  
Epoch 6/50  
141/141 [=====] - 219s 2s/step - loss: 3.4770 - accuracy: 0.1263 - va  
l_loss: 3.6056 - val_accuracy: 0.0798  
Epoch 7/50  
141/141 [=====] - 218s 2s/step - loss: 3.3913 - accuracy: 0.1354 - va  
l_loss: 3.4952 - val_accuracy: 0.1046  
Epoch 8/50  
141/141 [=====] - 208s 1s/step - loss: 3.3078 - accuracy: 0.1523 - va  
l_loss: 3.4411 - val_accuracy: 0.1117  
Epoch 9/50  
141/141 [=====] - 203s 1s/step - loss: 3.2560 - accuracy: 0.1602 - va  
l_loss: 3.3436 - val_accuracy: 0.1259  
Epoch 10/50  
141/141 [=====] - 218s 2s/step - loss: 3.1938 - accuracy: 0.1811 - va  
l_loss: 3.4025 - val_accuracy: 0.1383  
Epoch 11/50  
141/141 [=====] - 218s 2s/step - loss: 3.0845 - accuracy: 0.1968 - va  
l_loss: 3.7430 - val_accuracy: 0.1170  
Epoch 12/50  
141/141 [=====] - 218s 2s/step - loss: 3.0342 - accuracy: 0.2043 - va  
l_loss: 4.2336 - val_accuracy: 0.0443  
Epoch 13/50  
141/141 [=====] - 219s 2s/step - loss: 3.0092 - accuracy: 0.2090 - va  
l_loss: 3.8172 - val_accuracy: 0.0993  
Epoch 14/50  
141/141 [=====] - 218s 2s/step - loss: 2.9281 - accuracy: 0.2263 - va  
l_loss: 3.4043 - val_accuracy: 0.1436  
Epoch 15/50  
141/141 [=====] - 219s 2s/step - loss: 2.8614 - accuracy: 0.2389 - va  
l_loss: 3.3307 - val_accuracy: 0.1312  
Epoch 16/50  
141/141 [=====] - 216s 2s/step - loss: 2.7742 - accuracy: 0.2609 - va  
l_loss: 3.3919 - val_accuracy: 0.1401  
Epoch 17/50  
141/141 [=====] - 204s 1s/step - loss: 2.7013 - accuracy: 0.2777 - va  
l_loss: 3.2417 - val_accuracy: 0.1667  
Epoch 18/50  
141/141 [=====] - 213s 2s/step - loss: 2.6280 - accuracy: 0.2883 - va  
l_loss: 3.3154 - val_accuracy: 0.1418  
Epoch 19/50  
141/141 [=====] - 218s 2s/step - loss: 2.5648 - accuracy: 0.3147 - va  
l_loss: 3.2357 - val_accuracy: 0.1826  
Epoch 20/50  
141/141 [=====] - 218s 2s/step - loss: 2.4679 - accuracy: 0.3291 - va  
l_loss: 3.3177 - val_accuracy: 0.1649  
Epoch 21/50  
141/141 [=====] - 221s 2s/step - loss: 2.3797 - accuracy: 0.3473 - va  
l_loss: 3.6015 - val_accuracy: 0.1223  
Epoch 22/50  
141/141 [=====] - 218s 2s/step - loss: 2.2998 - accuracy: 0.3699 - va  
l_loss: 3.3148 - val_accuracy: 0.1809  
Epoch 23/50  
141/141 [=====] - 218s 2s/step - loss: 2.2309 - accuracy: 0.3870 - va  
l_loss: 3.6513 - val_accuracy: 0.1489  
Epoch 24/50  
141/141 [=====] - 220s 2s/step - loss: 2.1293 - accuracy: 0.4056 - va  
l_loss: 3.1309 - val_accuracy: 0.2128  
Epoch 25/50  
141/141 [=====] - 205s 1s/step - loss: 2.0382 - accuracy: 0.4333 - va  
l_loss: 3.2695 - val_accuracy: 0.1897  
Epoch 26/50  
141/141 [=====] - 207s 1s/step - loss: 1.9455 - accuracy: 0.4566 - va  
l_loss: 3.2975 - val_accuracy: 0.1986  
Epoch 27/50  
141/141 [=====] - 217s 2s/step - loss: 1.8819 - accuracy: 0.4843 - va  
l_loss: 3.5896 - val_accuracy: 0.1933
```

```
Epoch 28/50
141/141 [=====] - 218s 2s/step - loss: 1.7704 - accuracy: 0.4991 - val_loss: 3.6628 - val_accuracy: 0.1507
Epoch 29/50
141/141 [=====] - 218s 2s/step - loss: 1.6962 - accuracy: 0.5215 - val_loss: 3.1028 - val_accuracy: 0.2358
Epoch 30/50
141/141 [=====] - 218s 2s/step - loss: 1.6295 - accuracy: 0.5363 - val_loss: 3.4030 - val_accuracy: 0.2039
Epoch 31/50
141/141 [=====] - 218s 2s/step - loss: 1.5362 - accuracy: 0.5590 - val_loss: 3.7357 - val_accuracy: 0.1755
Epoch 32/50
141/141 [=====] - 218s 2s/step - loss: 1.4562 - accuracy: 0.5849 - val_loss: 3.1535 - val_accuracy: 0.2518
Epoch 33/50
141/141 [=====] - 209s 1s/step - loss: 1.3643 - accuracy: 0.6059 - val_loss: 3.1617 - val_accuracy: 0.2500
Epoch 34/50
141/141 [=====] - 204s 1s/step - loss: 1.2824 - accuracy: 0.6257 - val_loss: 3.2588 - val_accuracy: 0.2376
Epoch 35/50
141/141 [=====] - 217s 2s/step - loss: 1.1926 - accuracy: 0.6567 - val_loss: 3.2759 - val_accuracy: 0.2589
Epoch 36/50
141/141 [=====] - 218s 2s/step - loss: 1.1101 - accuracy: 0.6784 - val_loss: 3.2158 - val_accuracy: 0.2624
Epoch 37/50
141/141 [=====] - 218s 2s/step - loss: 1.0726 - accuracy: 0.6933 - val_loss: 3.3322 - val_accuracy: 0.2677
Epoch 38/50
141/141 [=====] - 218s 2s/step - loss: 1.0134 - accuracy: 0.7081 - val_loss: 3.2056 - val_accuracy: 0.2589
Epoch 39/50
141/141 [=====] - 217s 2s/step - loss: 0.9505 - accuracy: 0.7256 - val_loss: 3.7358 - val_accuracy: 0.2199
Epoch 40/50
141/141 [=====] - 218s 2s/step - loss: 0.8755 - accuracy: 0.7436 - val_loss: 3.3627 - val_accuracy: 0.2589
Epoch 41/50
141/141 [=====] - 214s 2s/step - loss: 0.8179 - accuracy: 0.7609 - val_loss: 3.5969 - val_accuracy: 0.2465
Epoch 42/50
141/141 [=====] - 203s 1s/step - loss: 0.7254 - accuracy: 0.7968 - val_loss: 3.4931 - val_accuracy: 0.2784
Epoch 43/50
141/141 [=====] - 210s 1s/step - loss: 0.7010 - accuracy: 0.7872 - val_loss: 3.3895 - val_accuracy: 0.2642
Epoch 44/50
141/141 [=====] - 217s 2s/step - loss: 0.6523 - accuracy: 0.8127 - val_loss: 4.2434 - val_accuracy: 0.2305
Epoch 45/50
141/141 [=====] - 218s 2s/step - loss: 0.6263 - accuracy: 0.8183 - val_loss: 3.5627 - val_accuracy: 0.2695
Epoch 46/50
141/141 [=====] - 218s 2s/step - loss: 0.5521 - accuracy: 0.8400 - val_loss: 3.5264 - val_accuracy: 0.2642
Epoch 47/50
141/141 [=====] - 218s 2s/step - loss: 0.5510 - accuracy: 0.8375 - val_loss: 3.6248 - val_accuracy: 0.2606
Epoch 48/50
141/141 [=====] - 217s 2s/step - loss: 0.4987 - accuracy: 0.8577 - val_loss: 3.5673 - val_accuracy: 0.2606
Epoch 49/50
141/141 [=====] - 219s 2s/step - loss: 0.4955 - accuracy: 0.8577 - val_loss: 3.7836 - val_accuracy: 0.2730
Epoch 50/50
```

```
141/141 [=====] - 207s 1s/step - loss: 0.4411 - accuracy: 0.8757 - val_loss: 3.6587 - val_accuracy: 0.2411
```



```
In [48]: # Evaluate the model on the testing set
score = model.evaluate(X_test, y_test, verbose=2)
print('Test Loss:', score[0])
print('Test accuracy:', score[1])
```

```
18/18 - 6s - loss: 3.6180 - accuracy: 0.2677 - 6s/epoch - 319ms/step
Test Loss: 3.6179747581481934
Test accuracy: 0.2677305042743683
```

```
In [49]: # Predict classes for the testing set
y_pred = model.predict(X_test)
y_pred = np.argmax(y_pred, axis=1)

# Convert probabilities to class labels
y_test = np.argmax(y_test, axis=1)

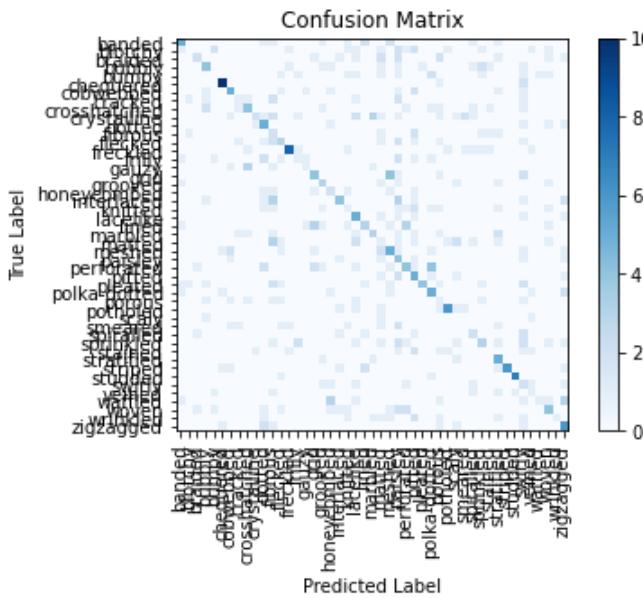
# Calculate precision, recall, and other metrics
print("Classification Report:")
print(classification_report(y_test, y_pred))

# Calculate and plot confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion Matrix')
plt.colorbar()
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.xticks(np.arange(len(class_names)), class_names, rotation=90)
plt.yticks(np.arange(len(class_names)), class_names)
plt.show()
```

18/18 [=====] - 6s 303ms/step

Classification Report:

	precision	recall	f1-score	support
0	0.42	0.33	0.37	15
1	0.00	0.00	0.00	16
2	0.40	0.14	0.21	14
3	0.33	0.27	0.30	15
4	0.00	0.00	0.00	12
5	0.71	0.77	0.74	13
6	0.38	0.36	0.37	14
7	0.20	0.10	0.13	10
8	0.36	0.29	0.32	14
9	0.50	0.11	0.17	19
10	0.26	0.50	0.34	10
11	0.08	0.22	0.12	9
12	0.23	0.23	0.23	13
13	0.67	0.47	0.55	17
14	0.17	0.09	0.12	11
15	0.25	0.14	0.18	7
16	0.33	0.33	0.33	12
17	0.27	0.33	0.30	9
18	0.18	0.25	0.21	8
19	0.40	0.25	0.31	16
20	0.10	0.08	0.09	12
21	0.25	0.50	0.33	10
22	0.33	0.20	0.25	15
23	0.33	0.25	0.29	12
24	0.05	0.07	0.06	15
25	0.28	0.31	0.29	16
26	0.11	0.40	0.17	10
27	0.29	0.24	0.26	17
28	0.19	0.45	0.26	11
29	0.80	0.25	0.38	16
30	0.29	0.42	0.34	12
31	0.06	0.11	0.08	9
32	0.60	0.50	0.55	12
33	0.00	0.00	0.00	3
34	0.00	0.00	0.00	9
35	0.20	0.25	0.22	8
36	0.27	0.18	0.21	17
37	0.00	0.00	0.00	9
38	0.26	0.62	0.37	8
39	1.00	0.46	0.63	13
40	0.88	0.78	0.82	9
41	0.07	0.22	0.11	9
42	0.14	0.17	0.15	6
43	0.50	0.08	0.14	12
44	0.29	0.25	0.27	16
45	1.00	0.09	0.17	11
46	0.27	0.46	0.34	13
accuracy			0.27	564
macro avg	0.31	0.27	0.26	564
weighted avg	0.34	0.27	0.27	564



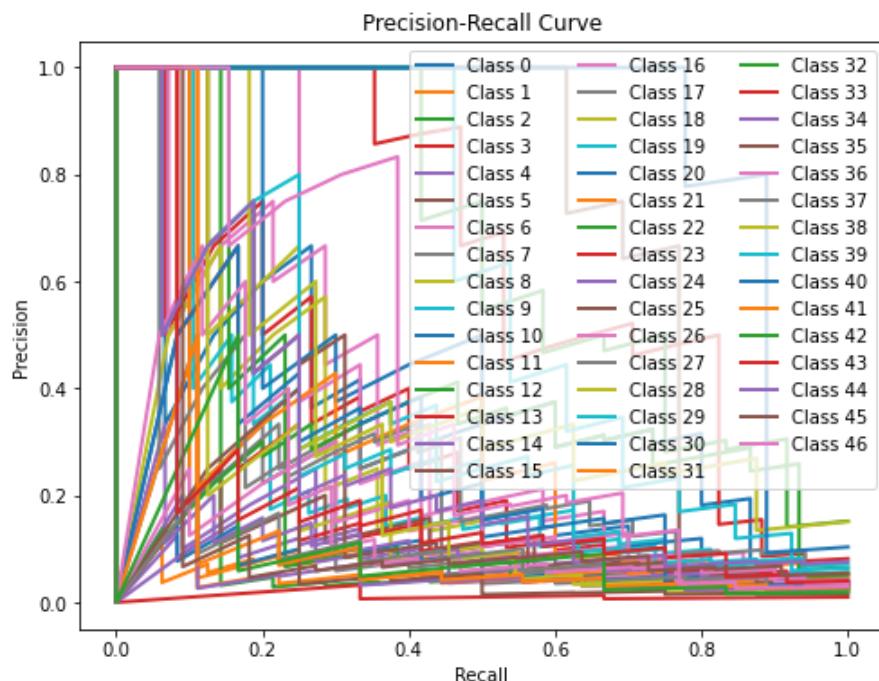
```
In [50]: # Obtain the predicted probabilities for each class
y_pred_prob = model.predict(X_test)

# Binarize the true labels
y_test_bin = label_binarize(y_test, classes=np.unique(y_test))

# Calculate precision and recall for each class
precision = dict()
recall = dict()
for i in range(num_classes):
    precision[i], recall[i], _ = precision_recall_curve(y_test_bin[:, i], y_pred_prob[:, i])

# Plot precision-recall curve for each class
plt.figure(figsize=(8, 6))
for i in range(num_classes):
    plt.plot(recall[i], precision[i], lw=2, label=f'Class {i}')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.legend(loc='upper right', bbox_to_anchor=(1.0, 1.0), ncol=3)
plt.show()
```

18/18 [=====] - 5s 295ms/step



Inception V4 Model

Inception V4 Model using Adam Optimizer

```
In [61]: # Load the pre-trained InceptionV4 model without the top layer
inceptionv4 = keras.applications.InceptionV3(
    include_top=False,
    weights='imagenet',
    input_shape=(224, 224, 3),
    pooling=None
)

# Add a new top layer for classification
x = layers.GlobalAveragePooling2D()(inceptionv4.output)
x = layers.Dense(1024, activation='relu')(x)
x = layers.Dropout(0.5)(x)
predictions = layers.Dense(num_classes, activation='softmax')(x)

# Define the new model
model = keras.Model(inputs=inceptionv4.input, outputs=predictions)

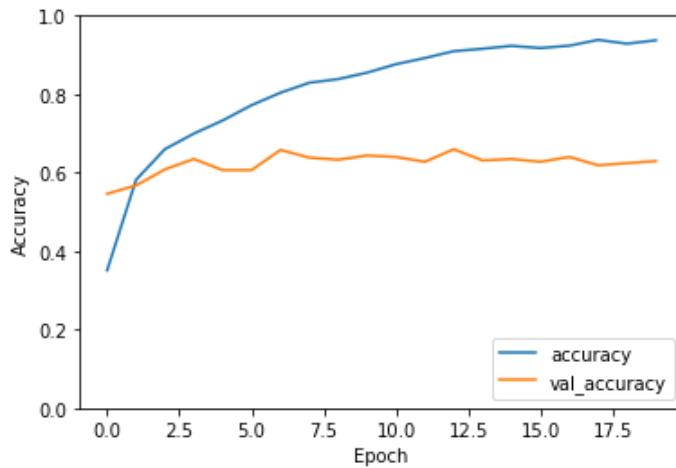
# Freeze the pre-trained layers
for layer in inceptionv4.layers:
    layer.trainable = False

# Compile the model
model.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

# Train the model
history = model.fit(
    X_train, y_train,
    epochs=20,
    batch_size=32,
    verbose=1,
    validation_data=(X_val, y_val)
)

plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0, 1])
plt.legend(loc='lower right')
plt.show()
```

```
Epoch 1/20
141/141 [=====] - 195s 1s/step - loss: 2.6797 - accuracy: 0.3511 - val_loss: 1.7677 - val_accuracy: 0.5461
Epoch 2/20
141/141 [=====] - 191s 1s/step - loss: 1.5153 - accuracy: 0.5816 - val_loss: 1.5624 - val_accuracy: 0.5674
Epoch 3/20
141/141 [=====] - 194s 1s/step - loss: 1.2100 - accuracy: 0.6598 - val_loss: 1.4754 - val_accuracy: 0.6082
Epoch 4/20
141/141 [=====] - 195s 1s/step - loss: 1.0368 - accuracy: 0.6995 - val_loss: 1.3463 - val_accuracy: 0.6348
Epoch 5/20
141/141 [=====] - 188s 1s/step - loss: 0.8762 - accuracy: 0.7329 - val_loss: 1.4532 - val_accuracy: 0.6064
Epoch 6/20
141/141 [=====] - 188s 1s/step - loss: 0.7406 - accuracy: 0.7719 - val_loss: 1.4581 - val_accuracy: 0.6064
Epoch 7/20
141/141 [=====] - 188s 1s/step - loss: 0.6212 - accuracy: 0.8036 - val_loss: 1.4069 - val_accuracy: 0.6578
Epoch 8/20
141/141 [=====] - 178s 1s/step - loss: 0.5506 - accuracy: 0.8291 - val_loss: 1.4487 - val_accuracy: 0.6383
Epoch 9/20
141/141 [=====] - 180s 1s/step - loss: 0.5058 - accuracy: 0.8384 - val_loss: 1.4913 - val_accuracy: 0.6330
Epoch 10/20
141/141 [=====] - 188s 1s/step - loss: 0.4467 - accuracy: 0.8548 - val_loss: 1.4454 - val_accuracy: 0.6436
Epoch 11/20
141/141 [=====] - 188s 1s/step - loss: 0.3833 - accuracy: 0.8763 - val_loss: 1.5249 - val_accuracy: 0.6401
Epoch 12/20
141/141 [=====] - 188s 1s/step - loss: 0.3378 - accuracy: 0.8921 - val_loss: 1.4971 - val_accuracy: 0.6277
Epoch 13/20
141/141 [=====] - 190s 1s/step - loss: 0.2922 - accuracy: 0.9096 - val_loss: 1.5887 - val_accuracy: 0.6596
Epoch 14/20
141/141 [=====] - 188s 1s/step - loss: 0.2661 - accuracy: 0.9158 - val_loss: 1.6276 - val_accuracy: 0.6312
Epoch 15/20
141/141 [=====] - 190s 1s/step - loss: 0.2389 - accuracy: 0.9231 - val_loss: 1.6050 - val_accuracy: 0.6348
Epoch 16/20
141/141 [=====] - 188s 1s/step - loss: 0.2436 - accuracy: 0.9176 - val_loss: 1.7046 - val_accuracy: 0.6277
Epoch 17/20
141/141 [=====] - 187s 1s/step - loss: 0.2356 - accuracy: 0.9235 - val_loss: 1.7589 - val_accuracy: 0.6401
Epoch 18/20
141/141 [=====] - 174s 1s/step - loss: 0.1963 - accuracy: 0.9384 - val_loss: 1.7652 - val_accuracy: 0.6188
Epoch 19/20
141/141 [=====] - 185s 1s/step - loss: 0.2159 - accuracy: 0.9286 - val_loss: 1.7723 - val_accuracy: 0.6241
Epoch 20/20
141/141 [=====] - 188s 1s/step - loss: 0.1978 - accuracy: 0.9373 - val_loss: 1.8616 - val_accuracy: 0.6294
```



```
In [62]: # Evaluate the model on the testing set
score = model.evaluate(X_test, y_test, verbose=2)
print('Test Loss:', score[0])
print('Test accuracy:', score[1])

18/18 - 21s - loss: 1.7414 - accuracy: 0.6401 - 21s/epoch - 1s/step
Test Loss: 1.7413880825042725
Test accuracy: 0.640070915222168
```

```
In [63]: # Predict classes for the testing set
y_pred = model.predict(X_test)
y_pred = np.argmax(y_pred, axis=1)

# Convert probabilities to class labels
y_test = np.argmax(y_test, axis=1)

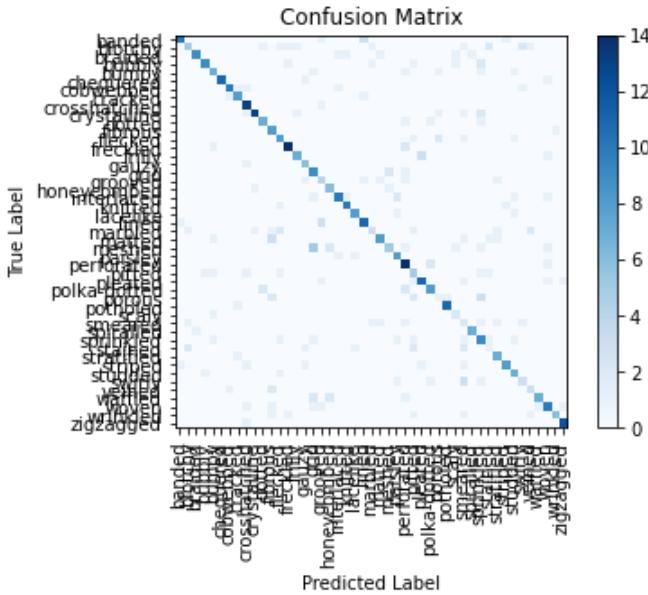
# Calculate precision, recall, and other metrics
print("Classification Report:")
print(classification_report(y_test, y_pred))

# Calculate and plot confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion Matrix')
plt.colorbar()
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.xticks(np.arange(len(class_names)), class_names, rotation=90)
plt.yticks(np.arange(len(class_names)), class_names)
plt.show()
```

18/18 [=====] - 23s 1s/step

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.67	0.74	15
1	0.56	0.31	0.40	16
2	0.90	0.64	0.75	14
3	0.82	0.60	0.69	15
4	0.58	0.58	0.58	12
5	1.00	0.85	0.92	13
6	0.83	0.71	0.77	14
7	0.73	0.80	0.76	10
8	0.65	0.93	0.76	14
9	0.87	0.68	0.76	19
10	0.70	0.70	0.70	10
11	0.44	0.89	0.59	9
12	0.53	0.62	0.57	13
13	0.82	0.82	0.82	17
14	0.78	0.64	0.70	11
15	0.60	0.86	0.71	7
16	0.45	0.75	0.56	12
17	0.36	0.44	0.40	9
18	0.60	0.75	0.67	8
19	0.77	0.62	0.69	16
20	0.83	0.83	0.83	12
21	0.80	0.80	0.80	10
22	0.69	0.73	0.71	15
23	0.50	0.25	0.33	12
24	0.73	0.53	0.62	15
25	0.62	0.31	0.42	16
26	0.62	0.80	0.70	10
27	0.56	0.82	0.67	17
28	0.33	0.45	0.38	11
29	0.73	0.69	0.71	16
30	0.67	0.67	0.67	12
31	0.20	0.11	0.14	9
32	0.92	0.92	0.92	12
33	0.40	0.67	0.50	3
34	0.13	0.22	0.17	9
35	0.88	0.88	0.88	8
36	0.41	0.53	0.46	17
37	0.17	0.11	0.13	9
38	0.64	0.88	0.74	8
39	0.80	0.62	0.70	13
40	0.58	0.78	0.67	9
41	0.43	0.33	0.38	9
42	0.43	0.50	0.46	6
43	0.88	0.58	0.70	12
44	0.62	0.62	0.62	16
45	0.86	0.55	0.67	11
46	0.86	0.92	0.89	13
accuracy			0.64	564
macro avg	0.64	0.64	0.63	564
weighted avg	0.66	0.64	0.64	564



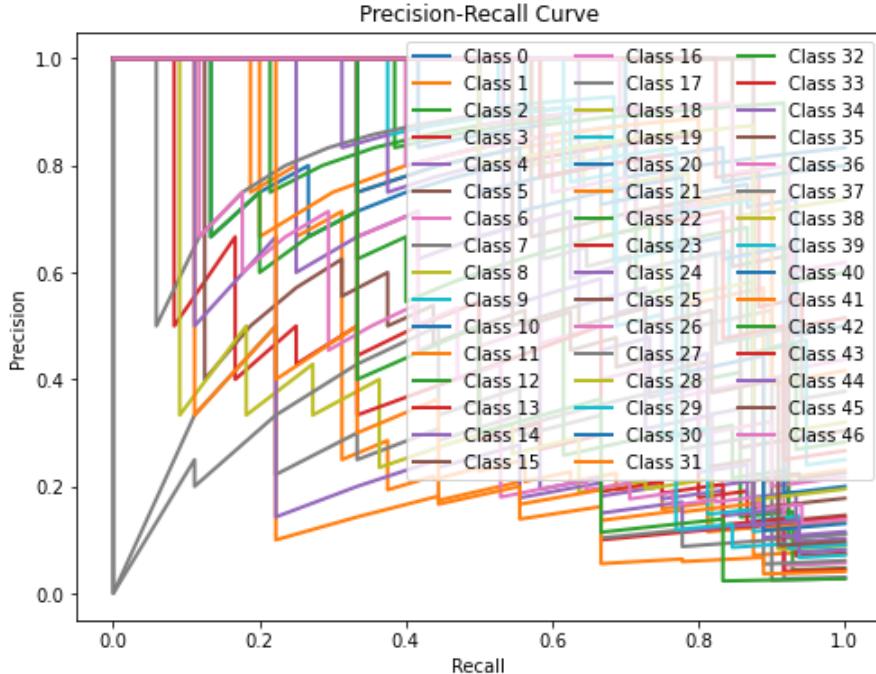
```
In [64]: # Obtain the predicted probabilities for each class
y_pred_prob = model.predict(X_test)

# Binarize the true labels
y_test_bin = label_binarize(y_test, classes=np.unique(y_test))

# Calculate precision and recall for each class
precision = dict()
recall = dict()
for i in range(num_classes):
    precision[i], recall[i], _ = precision_recall_curve(y_test_bin[:, i], y_pred_prob[:, i])

# Plot precision-recall curve for each class
plt.figure(figsize=(8, 6))
for i in range(num_classes):
    plt.plot(recall[i], precision[i], lw=2, label=f'Class {i}')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.legend(loc='upper right', bbox_to_anchor=(1.0, 1.0), ncol=3)
plt.show()
```

18/18 [=====] - 22s 1s/step



Inception V4 Model using SGD Optimizer with 20 Epochs

```
In [74]: # Load the pre-trained InceptionV4 model without the top layer
inceptionv4 = keras.applications.InceptionV3(
    include_top=False,
    weights='imagenet',
    input_shape=(224, 224, 3),
    pooling=None
)

# Add a new top layer for classification
x = layers.GlobalAveragePooling2D()(inceptionv4.output)
x = layers.Dense(1024, activation='relu')(x)
x = layers.Dropout(0.5)(x)
predictions = layers.Dense(num_classes, activation='softmax')(x)

# Define the new model
model = keras.Model(inputs=inceptionv4.input, outputs=predictions)

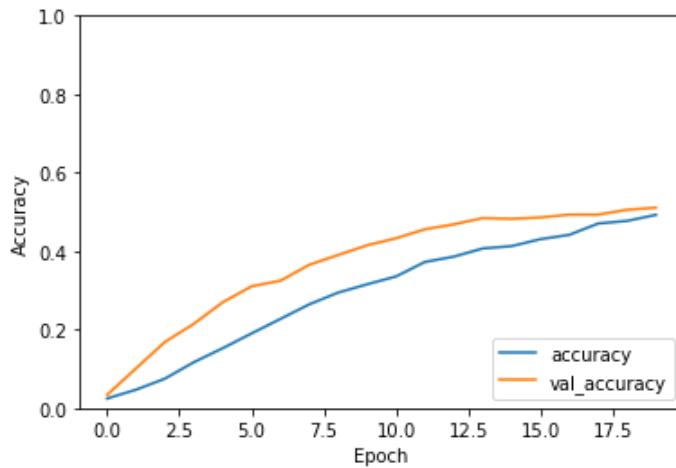
# Freeze the pre-trained layers
for layer in inceptionv4.layers:
    layer.trainable = False

# Compile the model
model.compile(
    optimizer=tf.optimizers.SGD(learning_rate=0.001),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

# Train the model
history = model.fit(
    X_train, y_train,
    epochs=20,
    batch_size=32,
    verbose=1,
    validation_data=(X_val, y_val)
)

plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0, 1])
plt.legend(loc='lower right')
plt.show()
```

```
Epoch 1/20
141/141 [=====] - 194s 1s/step - loss: 4.2191 - accuracy: 0.0244 - val_loss: 3.8667 - val_accuracy: 0.0337
Epoch 2/20
141/141 [=====] - 189s 1s/step - loss: 3.9424 - accuracy: 0.0468 - val_loss: 3.6944 - val_accuracy: 0.1011
Epoch 3/20
141/141 [=====] - 181s 1s/step - loss: 3.7428 - accuracy: 0.0749 - val_loss: 3.5489 - val_accuracy: 0.1684
Epoch 4/20
141/141 [=====] - 175s 1s/step - loss: 3.5939 - accuracy: 0.1168 - val_loss: 3.4218 - val_accuracy: 0.2145
Epoch 5/20
141/141 [=====] - 174s 1s/step - loss: 3.4284 - accuracy: 0.1525 - val_loss: 3.3053 - val_accuracy: 0.2695
Epoch 6/20
141/141 [=====] - 174s 1s/step - loss: 3.2982 - accuracy: 0.1902 - val_loss: 3.1970 - val_accuracy: 0.3103
Epoch 7/20
141/141 [=====] - 186s 1s/step - loss: 3.1808 - accuracy: 0.2272 - val_loss: 3.0935 - val_accuracy: 0.3245
Epoch 8/20
141/141 [=====] - 189s 1s/step - loss: 3.0475 - accuracy: 0.2644 - val_loss: 2.9918 - val_accuracy: 0.3652
Epoch 9/20
141/141 [=====] - 189s 1s/step - loss: 2.9452 - accuracy: 0.2943 - val_loss: 2.8995 - val_accuracy: 0.3901
Epoch 10/20
141/141 [=====] - 189s 1s/step - loss: 2.8348 - accuracy: 0.3154 - val_loss: 2.8088 - val_accuracy: 0.4149
Epoch 11/20
141/141 [=====] - 189s 1s/step - loss: 2.7562 - accuracy: 0.3353 - val_loss: 2.7235 - val_accuracy: 0.4326
Epoch 12/20
141/141 [=====] - 190s 1s/step - loss: 2.6477 - accuracy: 0.3723 - val_loss: 2.6470 - val_accuracy: 0.4557
Epoch 13/20
141/141 [=====] - 188s 1s/step - loss: 2.5714 - accuracy: 0.3859 - val_loss: 2.5732 - val_accuracy: 0.4681
Epoch 14/20
141/141 [=====] - 189s 1s/step - loss: 2.4900 - accuracy: 0.4067 - val_loss: 2.5009 - val_accuracy: 0.4840
Epoch 15/20
141/141 [=====] - 186s 1s/step - loss: 2.4252 - accuracy: 0.4127 - val_loss: 2.4344 - val_accuracy: 0.4823
Epoch 16/20
141/141 [=====] - 178s 1s/step - loss: 2.3608 - accuracy: 0.4304 - val_loss: 2.3747 - val_accuracy: 0.4858
Epoch 17/20
141/141 [=====] - 243s 2s/step - loss: 2.2846 - accuracy: 0.4415 - val_loss: 2.3161 - val_accuracy: 0.4929
Epoch 18/20
141/141 [=====] - 26649s 190s/step - loss: 2.2088 - accuracy: 0.4703 - val_loss: 2.2647 - val_accuracy: 0.4929
Epoch 19/20
141/141 [=====] - 290s 2s/step - loss: 2.1726 - accuracy: 0.4770 - val_loss: 2.2161 - val_accuracy: 0.5053
Epoch 20/20
141/141 [=====] - 295s 2s/step - loss: 2.0925 - accuracy: 0.4925 - val_loss: 2.1675 - val_accuracy: 0.5106
```



```
In [75]: # Evaluate the model on the testing set
score = model.evaluate(X_test, y_test, verbose=2)
print('Test Loss:', score[0])
print('Test accuracy:', score[1])

18/18 - 32s - loss: 2.0861 - accuracy: 0.5177 - 32s/epoch - 2s/step
Test Loss: 2.0860631465911865
Test accuracy: 0.5177304744720459
```

```
In [76]: # Predict classes for the testing set
y_pred = model.predict(X_test)
y_pred = np.argmax(y_pred, axis=1)

# Convert probabilities to class labels
y_test = np.argmax(y_test, axis=1)

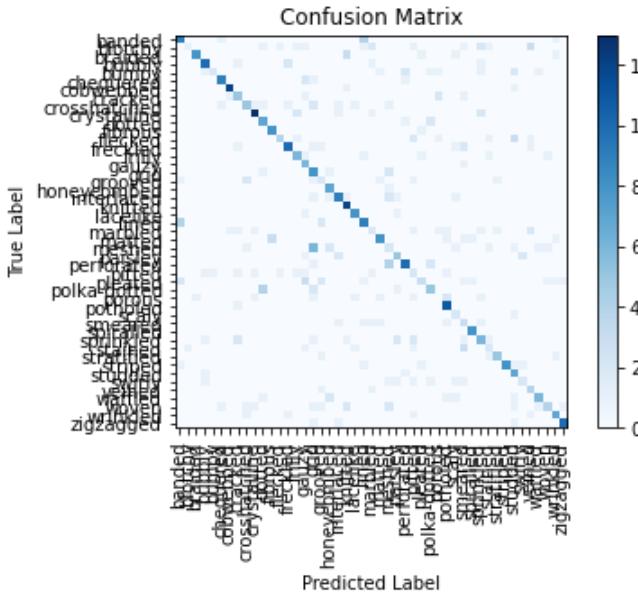
# Calculate precision, recall, and other metrics
print("Classification Report:")
print(classification_report(y_test, y_pred))

# Calculate and plot confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion Matrix')
plt.colorbar()
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.xticks(np.arange(len(class_names)), class_names, rotation=90)
plt.yticks(np.arange(len(class_names)), class_names)
plt.show()
```

18/18 [=====] - 39s 2s/step

Classification Report:

	precision	recall	f1-score	support
0	0.53	0.60	0.56	15
1	0.50	0.06	0.11	16
2	0.89	0.57	0.70	14
3	0.67	0.67	0.67	15
4	0.67	0.17	0.27	12
5	0.75	0.69	0.72	13
6	0.86	0.86	0.86	14
7	0.42	0.50	0.45	10
8	0.50	0.36	0.42	14
9	0.72	0.68	0.70	19
10	0.47	0.70	0.56	10
11	0.73	0.89	0.80	9
12	0.67	0.31	0.42	13
13	0.71	0.59	0.65	17
14	0.55	0.55	0.55	11
15	0.33	0.86	0.48	7
16	0.30	0.67	0.41	12
17	0.20	0.22	0.21	9
18	0.54	0.88	0.67	8
19	0.60	0.56	0.58	16
20	0.60	1.00	0.75	12
21	0.73	0.80	0.76	10
22	0.50	0.60	0.55	15
23	0.22	0.17	0.19	12
24	0.67	0.53	0.59	15
25	0.13	0.12	0.13	16
26	0.45	0.50	0.48	10
27	0.67	0.59	0.62	17
28	0.15	0.18	0.17	11
29	0.43	0.19	0.26	16
30	0.45	0.42	0.43	12
31	0.00	0.00	0.00	9
32	0.61	0.92	0.73	12
33	0.25	0.67	0.36	3
34	0.20	0.33	0.25	9
35	0.80	1.00	0.89	8
36	0.40	0.35	0.38	17
37	0.25	0.22	0.24	9
38	0.71	0.62	0.67	8
39	1.00	0.62	0.76	13
40	0.37	0.78	0.50	9
41	0.40	0.22	0.29	9
42	0.33	0.67	0.44	6
43	0.67	0.50	0.57	12
44	0.57	0.25	0.35	16
45	0.58	0.64	0.61	11
46	0.83	0.77	0.80	13
accuracy			0.52	564
macro avg	0.52	0.53	0.50	564
weighted avg	0.54	0.52	0.51	564



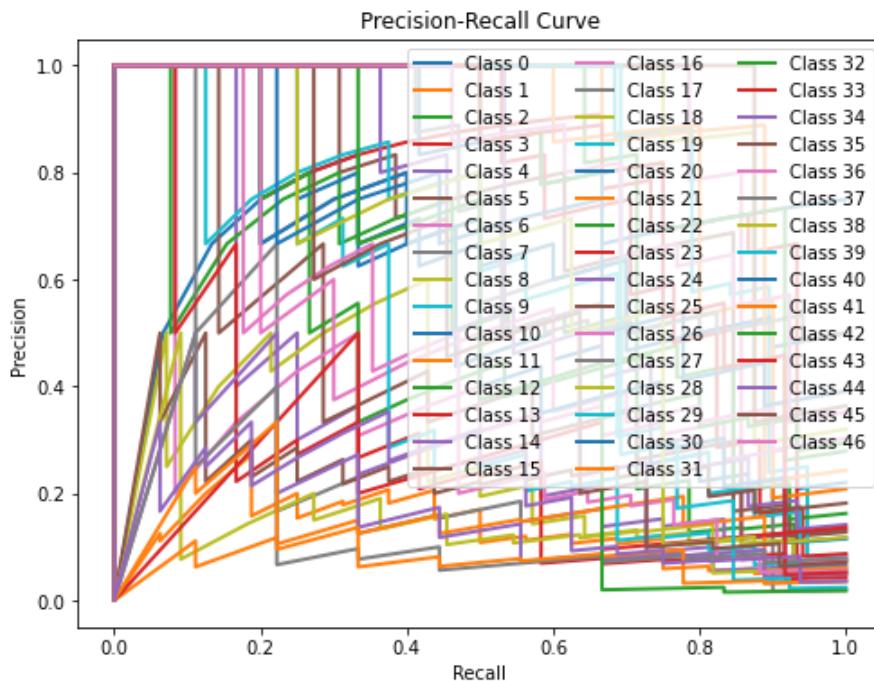
```
In [77]: # Obtain the predicted probabilities for each class
y_pred_prob = model.predict(X_test)

# Binarize the true labels
y_test_bin = label_binarize(y_test, classes=np.unique(y_test))

# Calculate precision and recall for each class
precision = dict()
recall = dict()
for i in range(num_classes):
    precision[i], recall[i], _ = precision_recall_curve(y_test_bin[:, i], y_pred_prob[:, i])

# Plot precision-recall curve for each class
plt.figure(figsize=(8, 6))
for i in range(num_classes):
    plt.plot(recall[i], precision[i], lw=2, label=f'Class {i}')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.legend(loc='upper right', bbox_to_anchor=(1.0, 1.0), ncol=3)
plt.show()
```

18/18 [=====] - 32s 2s/step



Inception V4 Model using SGD Optimizer with 50 Epochs

```
In [94]: # Load the pre-trained InceptionV4 model without the top layer
inceptionv4 = keras.applications.InceptionV3(
    include_top=False,
    weights='imagenet',
    input_shape=(224, 224, 3),
    pooling=None
)

# Add a new top layer for classification
x = layers.GlobalAveragePooling2D()(inceptionv4.output)
x = layers.Dense(1024, activation='relu')(x)
x = layers.Dropout(0.5)(x)
predictions = layers.Dense(num_classes, activation='softmax')(x)

# Define the new model
model = keras.Model(inputs=inceptionv4.input, outputs=predictions)

# Freeze the pre-trained layers
for layer in inceptionv4.layers:
    layer.trainable = False

# Compile the model
model.compile(
    optimizer=tf.optimizers.SGD(learning_rate=0.001),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

# Train the model
history = model.fit(
    X_train, y_train,
    epochs=50,
    batch_size=32,
    verbose=1,
    validation_data=(X_val, y_val)
)

plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0, 1])
plt.legend(loc='lower right')
plt.show()
```

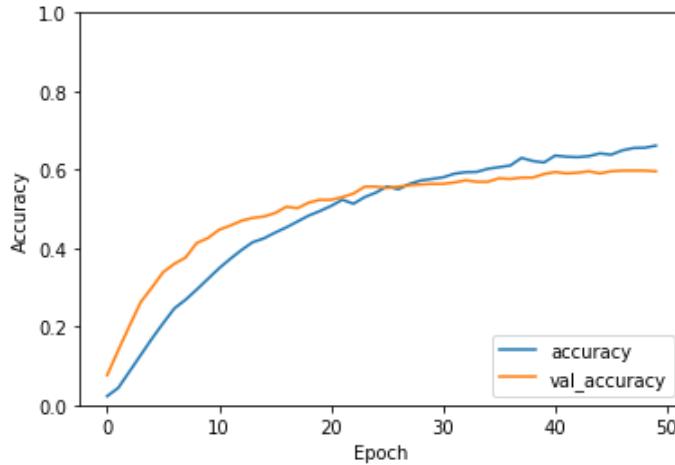
```
Epoch 1/50
141/141 [=====] - 308s 2s/step - loss: 4.2053 - accuracy: 0.0226 - val_loss: 3.7669 - val_accuracy: 0.0762
Epoch 2/50
141/141 [=====] - 273s 2s/step - loss: 3.8973 - accuracy: 0.0439 - val_loss: 3.6132 - val_accuracy: 0.1401
Epoch 3/50
141/141 [=====] - 248s 2s/step - loss: 3.7045 - accuracy: 0.0858 - val_loss: 3.4800 - val_accuracy: 0.2021
Epoch 4/50
141/141 [=====] - 189s 1s/step - loss: 3.5298 - accuracy: 0.1274 - val_loss: 3.3581 - val_accuracy: 0.2624
Epoch 5/50
141/141 [=====] - 207s 1s/step - loss: 3.3801 - accuracy: 0.1693 - val_loss: 3.2417 - val_accuracy: 0.2996
Epoch 6/50
141/141 [=====] - 202s 1s/step - loss: 3.2572 - accuracy: 0.2086 - val_loss: 3.1301 - val_accuracy: 0.3387
Epoch 7/50
141/141 [=====] - 210s 1s/step - loss: 3.1176 - accuracy: 0.2462 - val_loss: 3.0272 - val_accuracy: 0.3599
Epoch 8/50
141/141 [=====] - 204s 1s/step - loss: 3.0153 - accuracy: 0.2684 - val_loss: 2.9276 - val_accuracy: 0.3759
Epoch 9/50
141/141 [=====] - 200s 1s/step - loss: 2.9054 - accuracy: 0.2943 - val_loss: 2.8334 - val_accuracy: 0.4131
Epoch 10/50
141/141 [=====] - 196s 1s/step - loss: 2.7934 - accuracy: 0.3216 - val_loss: 2.7439 - val_accuracy: 0.4255
Epoch 11/50
141/141 [=====] - 199s 1s/step - loss: 2.7249 - accuracy: 0.3486 - val_loss: 2.6631 - val_accuracy: 0.4468
Epoch 12/50
141/141 [=====] - 212s 2s/step - loss: 2.6170 - accuracy: 0.3728 - val_loss: 2.5830 - val_accuracy: 0.4574
Epoch 13/50
141/141 [=====] - 194s 1s/step - loss: 2.5321 - accuracy: 0.3952 - val_loss: 2.5070 - val_accuracy: 0.4699
Epoch 14/50
141/141 [=====] - 192s 1s/step - loss: 2.4507 - accuracy: 0.4151 - val_loss: 2.4347 - val_accuracy: 0.4770
Epoch 15/50
141/141 [=====] - 180s 1s/step - loss: 2.3984 - accuracy: 0.4249 - val_loss: 2.3754 - val_accuracy: 0.4805
Epoch 16/50
141/141 [=====] - 188s 1s/step - loss: 2.3223 - accuracy: 0.4395 - val_loss: 2.3174 - val_accuracy: 0.4894
Epoch 17/50
141/141 [=====] - 183s 1s/step - loss: 2.2603 - accuracy: 0.4528 - val_loss: 2.2599 - val_accuracy: 0.5053
Epoch 18/50
141/141 [=====] - 177s 1s/step - loss: 2.1867 - accuracy: 0.4679 - val_loss: 2.2037 - val_accuracy: 0.5018
Epoch 19/50
141/141 [=====] - 178s 1s/step - loss: 2.1221 - accuracy: 0.4829 - val_loss: 2.1542 - val_accuracy: 0.5160
Epoch 20/50
141/141 [=====] - 187s 1s/step - loss: 2.0734 - accuracy: 0.4945 - val_loss: 2.1104 - val_accuracy: 0.5230
Epoch 21/50
141/141 [=====] - 190s 1s/step - loss: 2.0360 - accuracy: 0.5075 - val_loss: 2.0699 - val_accuracy: 0.5230
Epoch 22/50
141/141 [=====] - 196s 1s/step - loss: 1.9736 - accuracy: 0.5233 - val_loss: 2.0335 - val_accuracy: 0.5301
Epoch 23/50
141/141 [=====] - 192s 1s/step - loss: 1.9474 - accuracy: 0.5131 - val_loss: 2.0000 - val_accuracy: 0.5301
```

l_loss: 1.9941 - val_accuracy: 0.5390
Epoch 24/50
141/141 [=====] - 192s 1s/step - loss: 1.9078 - accuracy: 0.5301 - va
l_loss: 1.9589 - val_accuracy: 0.5567
Epoch 25/50
141/141 [=====] - 191s 1s/step - loss: 1.8622 - accuracy: 0.5414 - va
l_loss: 1.9273 - val_accuracy: 0.5567
Epoch 26/50
141/141 [=====] - 7267s 52s/step - loss: 1.8251 - accuracy: 0.5572 -
val_loss: 1.9022 - val_accuracy: 0.5532
Epoch 27/50
141/141 [=====] - 254s 2s/step - loss: 1.7921 - accuracy: 0.5499 - va
l_loss: 1.8730 - val_accuracy: 0.5567
Epoch 28/50
141/141 [=====] - 261s 2s/step - loss: 1.7471 - accuracy: 0.5636 - va
l_loss: 1.8494 - val_accuracy: 0.5603
Epoch 29/50
141/141 [=====] - 252s 2s/step - loss: 1.7119 - accuracy: 0.5723 - va
l_loss: 1.8235 - val_accuracy: 0.5621
Epoch 30/50
141/141 [=====] - 244s 2s/step - loss: 1.6806 - accuracy: 0.5762 - va
l_loss: 1.8010 - val_accuracy: 0.5638
Epoch 31/50
141/141 [=====] - 253s 2s/step - loss: 1.6706 - accuracy: 0.5805 - va
l_loss: 1.7796 - val_accuracy: 0.5638
Epoch 32/50
141/141 [=====] - 260s 2s/step - loss: 1.6390 - accuracy: 0.5895 - va
l_loss: 1.7600 - val_accuracy: 0.5674
Epoch 33/50
141/141 [=====] - 285s 2s/step - loss: 1.6114 - accuracy: 0.5935 - va
l_loss: 1.7415 - val_accuracy: 0.5727
Epoch 34/50
141/141 [=====] - 274s 2s/step - loss: 1.6093 - accuracy: 0.5942 - va
l_loss: 1.7247 - val_accuracy: 0.5691
Epoch 35/50
141/141 [=====] - 259s 2s/step - loss: 1.5563 - accuracy: 0.6022 - va
l_loss: 1.7094 - val_accuracy: 0.5691
Epoch 36/50
141/141 [=====] - 277s 2s/step - loss: 1.5416 - accuracy: 0.6064 - va
l_loss: 1.6944 - val_accuracy: 0.5780
Epoch 37/50
141/141 [=====] - 345s 2s/step - loss: 1.5320 - accuracy: 0.6106 - va
l_loss: 1.6802 - val_accuracy: 0.5762
Epoch 38/50
141/141 [=====] - 194s 1s/step - loss: 1.4944 - accuracy: 0.6299 - va
l_loss: 1.6654 - val_accuracy: 0.5798
Epoch 39/50
141/141 [=====] - 199s 1s/step - loss: 1.4747 - accuracy: 0.6221 - va
l_loss: 1.6516 - val_accuracy: 0.5798
Epoch 40/50
141/141 [=====] - 203s 1s/step - loss: 1.4760 - accuracy: 0.6181 - va
l_loss: 1.6415 - val_accuracy: 0.5887
Epoch 41/50
141/141 [=====] - 204s 1s/step - loss: 1.4365 - accuracy: 0.6359 - va
l_loss: 1.6288 - val_accuracy: 0.5940
Epoch 42/50
141/141 [=====] - 197s 1s/step - loss: 1.4207 - accuracy: 0.6330 - va
l_loss: 1.6179 - val_accuracy: 0.5904
Epoch 43/50
141/141 [=====] - 192s 1s/step - loss: 1.4128 - accuracy: 0.6319 - va
l_loss: 1.6073 - val_accuracy: 0.5922
Epoch 44/50
141/141 [=====] - 182s 1s/step - loss: 1.3996 - accuracy: 0.6343 - va
l_loss: 1.5994 - val_accuracy: 0.5957
Epoch 45/50
141/141 [=====] - 178s 1s/step - loss: 1.3721 - accuracy: 0.6414 - va
l_loss: 1.5907 - val_accuracy: 0.5904
Epoch 46/50

```

141/141 [=====] - 182s 1s/step - loss: 1.3627 - accuracy: 0.6381 - va
l_loss: 1.5764 - val_accuracy: 0.5957
Epoch 47/50
141/141 [=====] - 182s 1s/step - loss: 1.3490 - accuracy: 0.6492 - va
l_loss: 1.5667 - val_accuracy: 0.5975
Epoch 48/50
141/141 [=====] - 192s 1s/step - loss: 1.3264 - accuracy: 0.6551 - va
l_loss: 1.5623 - val_accuracy: 0.5975
Epoch 49/50
141/141 [=====] - 202s 1s/step - loss: 1.3078 - accuracy: 0.6558 - va
l_loss: 1.5541 - val_accuracy: 0.5975
Epoch 50/50
141/141 [=====] - 190s 1s/step - loss: 1.2964 - accuracy: 0.6611 - va
l_loss: 1.5475 - val_accuracy: 0.5957

```



```
In [96]: # Evaluate the model on the testing set
score = model.evaluate(X_test, y_test, verbose=2)
print('Test Loss:', score[0])
print('Test accuracy:', score[1])
```

```
18/18 - 21s - loss: 1.5318 - accuracy: 0.6064 - 21s/epoch - 1s/step
Test Loss: 1.5318270921707153
Test accuracy: 0.6063829660415649
```

```
In [97]: # Predict classes for the testing set
y_pred = model.predict(X_test)
y_pred = np.argmax(y_pred, axis=1)

# Convert probabilities to class labels
y_test = np.argmax(y_test, axis=1)

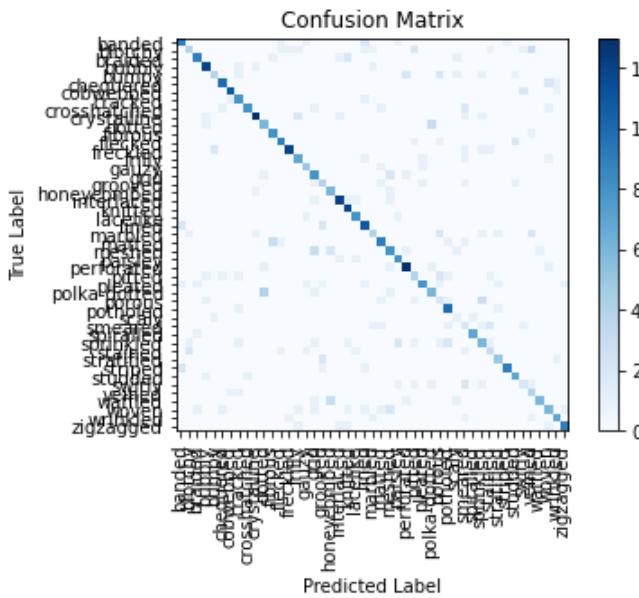
# Calculate precision, recall, and other metrics
print("Classification Report:")
print(classification_report(y_test, y_pred))

# Calculate and plot confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion Matrix')
plt.colorbar()
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.xticks(np.arange(len(class_names)), class_names, rotation=90)
plt.yticks(np.arange(len(class_names)), class_names)
plt.show()
```

18/18 [=====] - 24s 1s/step

Classification Report:

	precision	recall	f1-score	support
0	0.64	0.60	0.62	15
1	0.50	0.25	0.33	16
2	0.75	0.64	0.69	14
3	0.71	0.80	0.75	15
4	0.57	0.33	0.42	12
5	0.77	0.77	0.77	13
6	0.85	0.79	0.81	14
7	0.57	0.80	0.67	10
8	0.89	0.57	0.70	14
9	0.81	0.68	0.74	19
10	0.50	0.60	0.55	10
11	0.67	0.89	0.76	9
12	0.69	0.69	0.69	13
13	0.86	0.71	0.77	17
14	0.58	0.64	0.61	11
15	0.56	0.71	0.63	7
16	0.44	0.67	0.53	12
17	0.40	0.44	0.42	9
18	0.55	0.75	0.63	8
19	0.75	0.75	0.75	16
20	0.67	1.00	0.80	12
21	0.73	0.80	0.76	10
22	0.58	0.73	0.65	15
23	0.50	0.33	0.40	12
24	0.60	0.60	0.60	15
25	0.50	0.50	0.50	16
26	0.62	0.80	0.70	10
27	0.72	0.76	0.74	17
28	0.40	0.36	0.38	11
29	0.73	0.50	0.59	16
30	0.55	0.50	0.52	12
31	0.18	0.22	0.20	9
32	0.62	0.83	0.71	12
33	0.20	0.33	0.25	3
34	0.15	0.22	0.18	9
35	0.88	0.88	0.88	8
36	0.40	0.35	0.38	17
37	0.33	0.33	0.33	9
38	0.62	0.62	0.62	8
39	0.90	0.69	0.78	13
40	0.64	0.78	0.70	9
41	0.33	0.22	0.27	9
42	0.33	0.67	0.44	6
43	0.86	0.50	0.63	12
44	0.50	0.38	0.43	16
45	0.70	0.64	0.67	11
46	0.82	0.69	0.75	13
accuracy			0.61	564
macro avg	0.60	0.60	0.59	564
weighted avg	0.62	0.61	0.60	564



```
In [98]: # Obtain the predicted probabilities for each class
y_pred_prob = model.predict(X_test)

# Binarize the true labels
y_test_bin = label_binarize(y_test, classes=np.unique(y_test))

# Calculate precision and recall for each class
precision = dict()
recall = dict()
for i in range(num_classes):
    precision[i], recall[i], _ = precision_recall_curve(y_test_bin[:, i], y_pred_prob[:, i])

# Plot precision-recall curve for each class
plt.figure(figsize=(8, 6))
for i in range(num_classes):
    plt.plot(recall[i], precision[i], lw=2, label=f'Class {i}')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.legend(loc='upper right', bbox_to_anchor=(1.0, 1.0), ncol=3)
plt.show()
```

18/18 [=====] - 21s 1s/step

